

# IBM Cloud Private Performance and Sizing



IBM Cloud

# ICP Performance Analysis

Performance is analyzed based upon 3 representative application architectures

General performance analysis

## Sizing your Deployment

## Tuning your Cluster

# **ICP Performance Analysis**

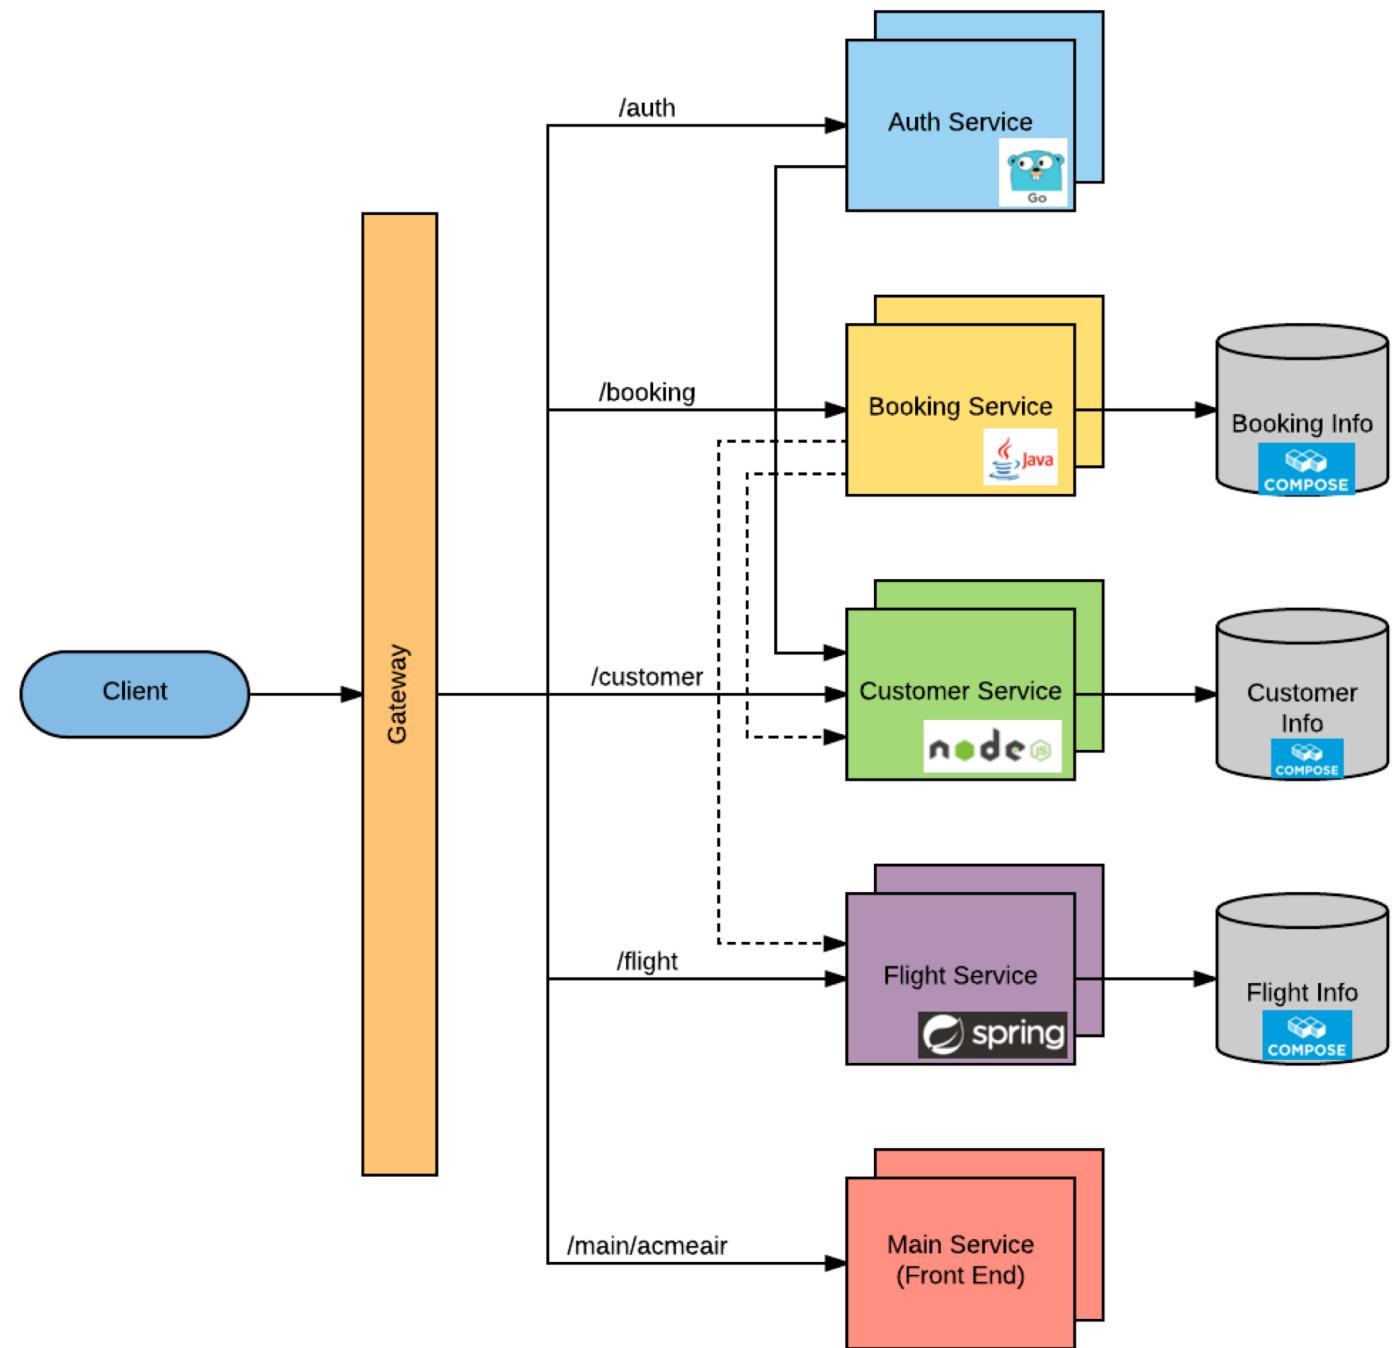
## Analyzing Performance with Applications

# The Polyglot

We will consider a Polyglot application referred to as Acme Air.

This application utilizes multiple programming frameworks.

The intent is to use a general case to evaluate performance and scalability of IBM Cloud Private.



# Online Banking

A second more complicated application is used to simulate Online Banking

Simulation of retail online banking (UI, security and services)

Traffic is encrypted and users are authenticated

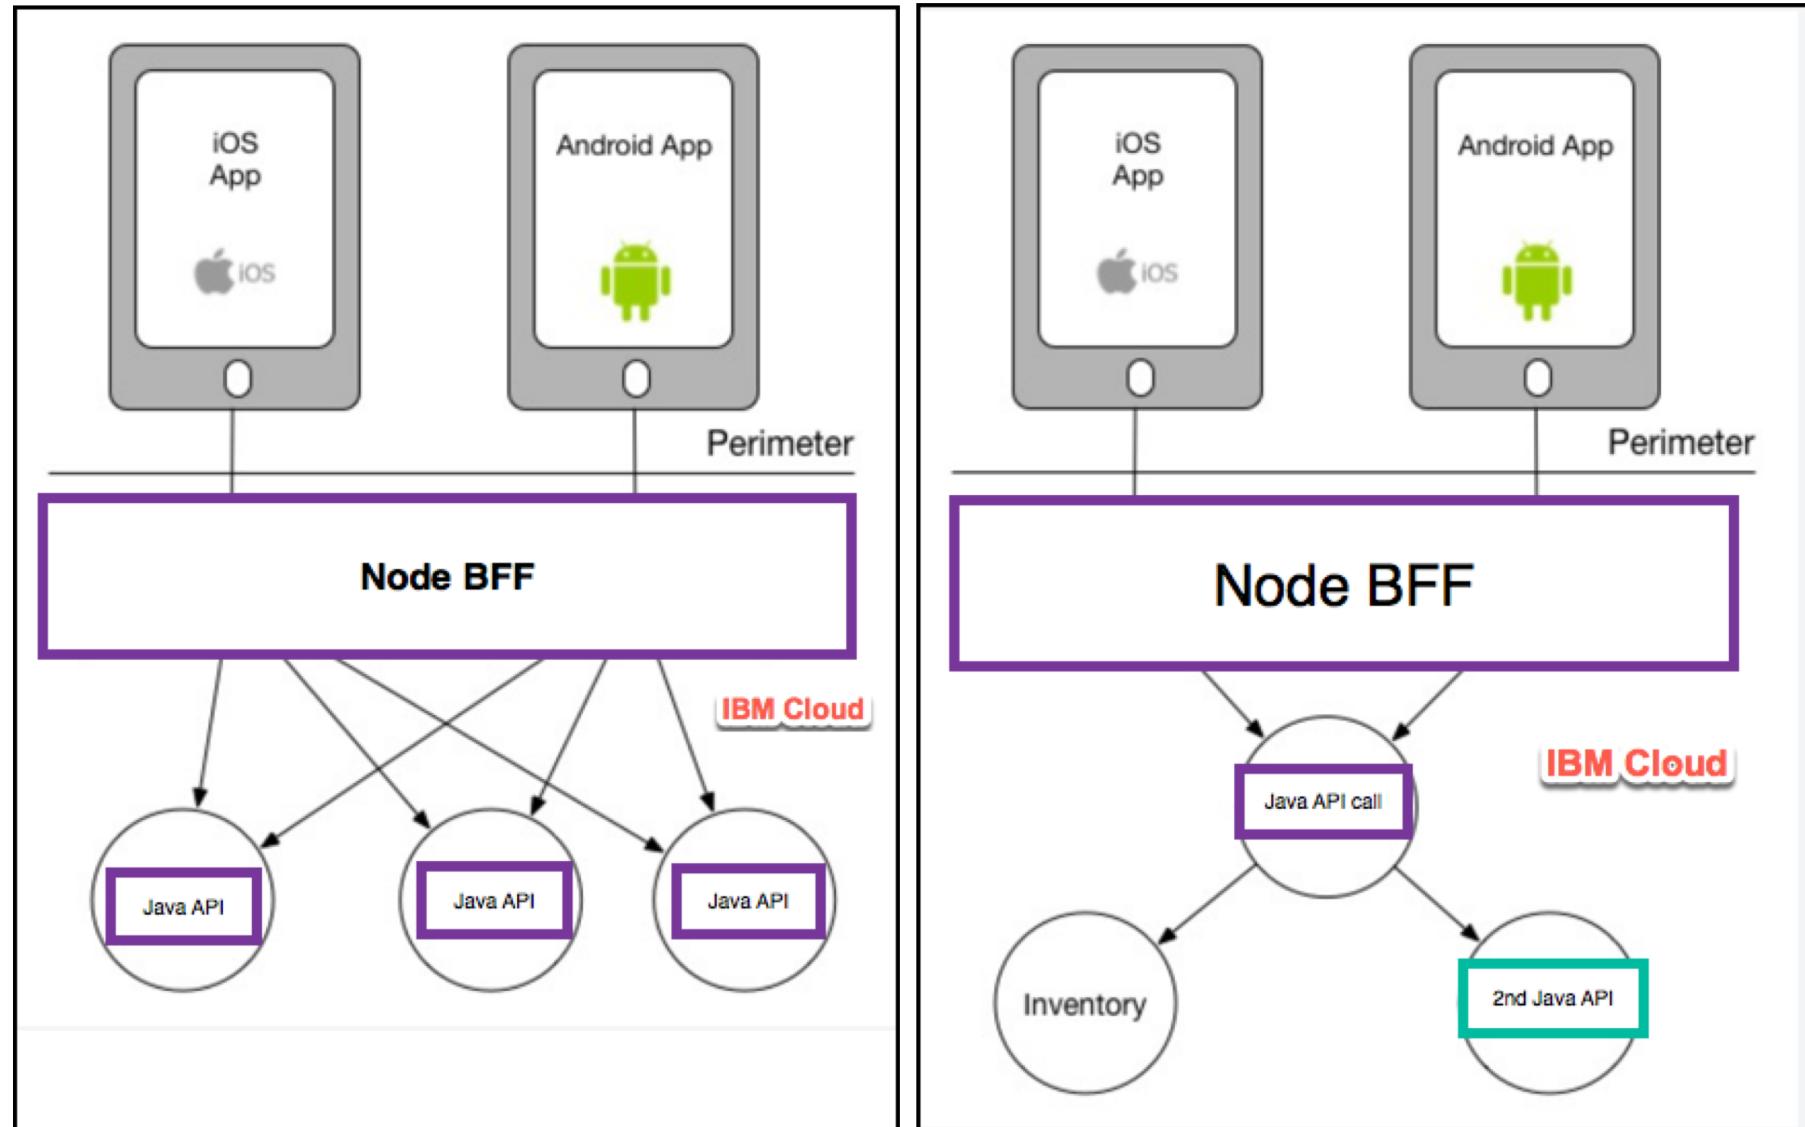
Backend services interaction has been simulated via Stub Application

Account Summary Page is developed using Angular JS

# Healthcare Microservices

Typical  
Microservices  
application

Architecture  
comparing 2 and  
3 levels

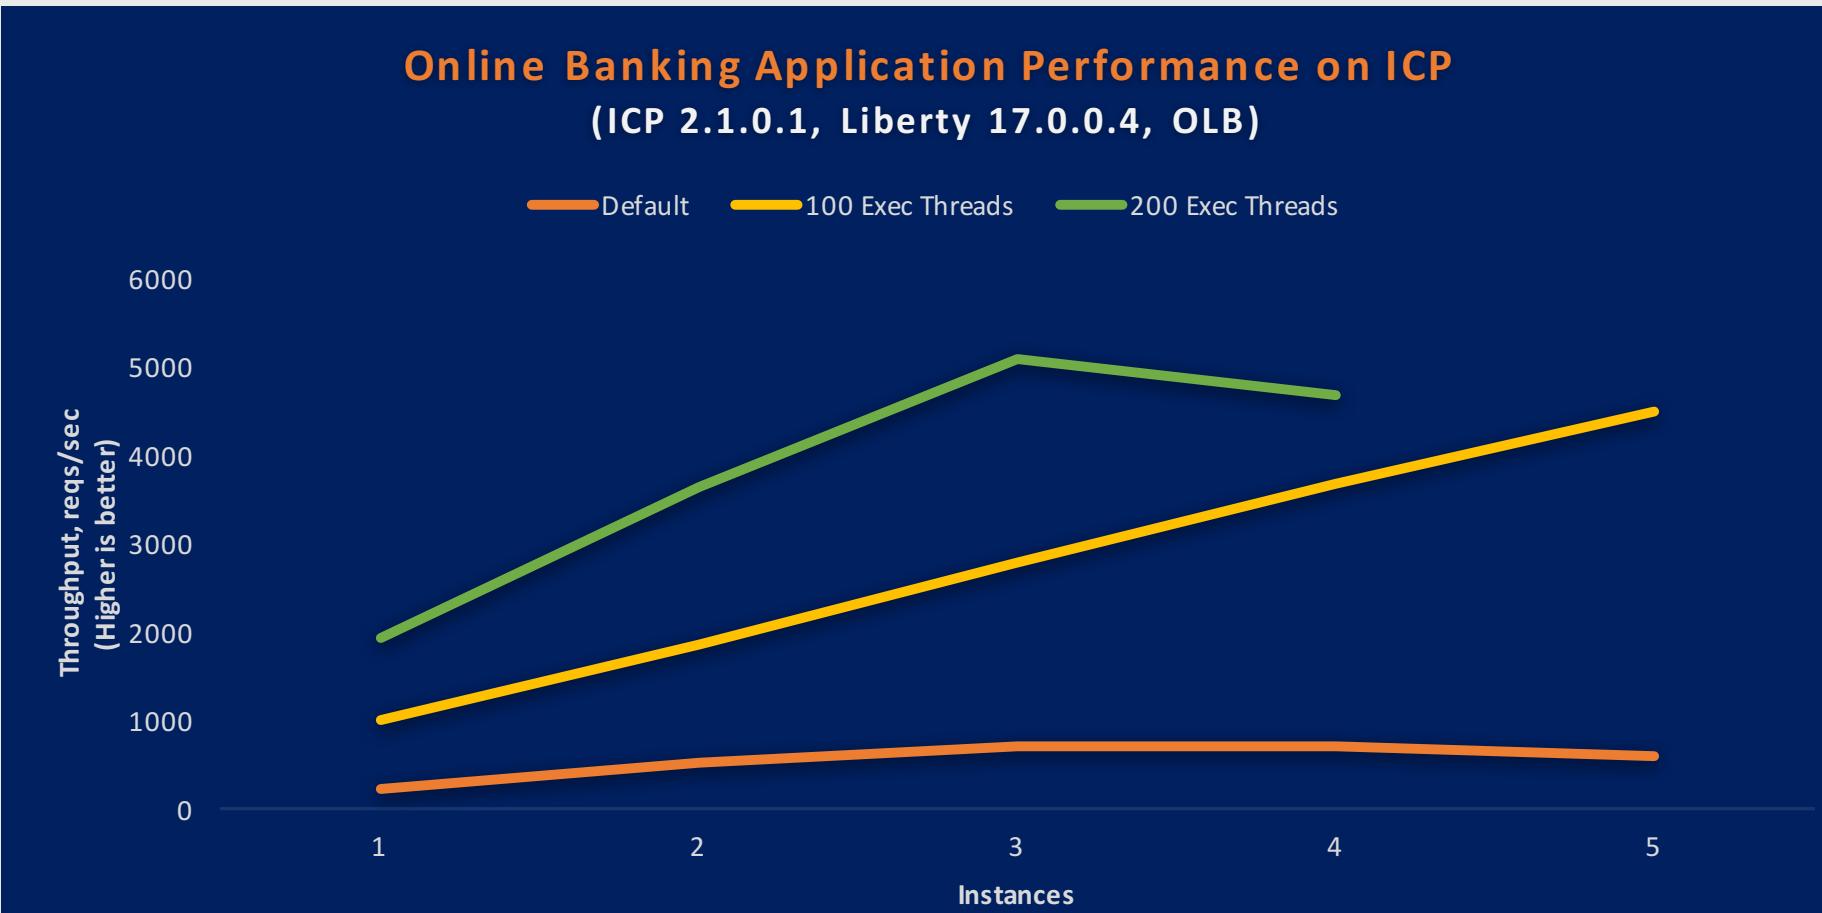


# Online Banking Performance on ICP

This type of Java microservice application scales nearly linearly  
Scaling this application, you have massive parallel workload coming in  
The runtime (in this case Liberty) must have enough executable threads to manage  
Cannot get full use of ICP resources without tuning  
Able to saturate by increasing executable threads

The default (orange) line shows only using horizontal scaling does not fully leverage the cluster resources

Understanding the application and its run time shows increasing the Liberty executor thread pool size is crucial for this Java application



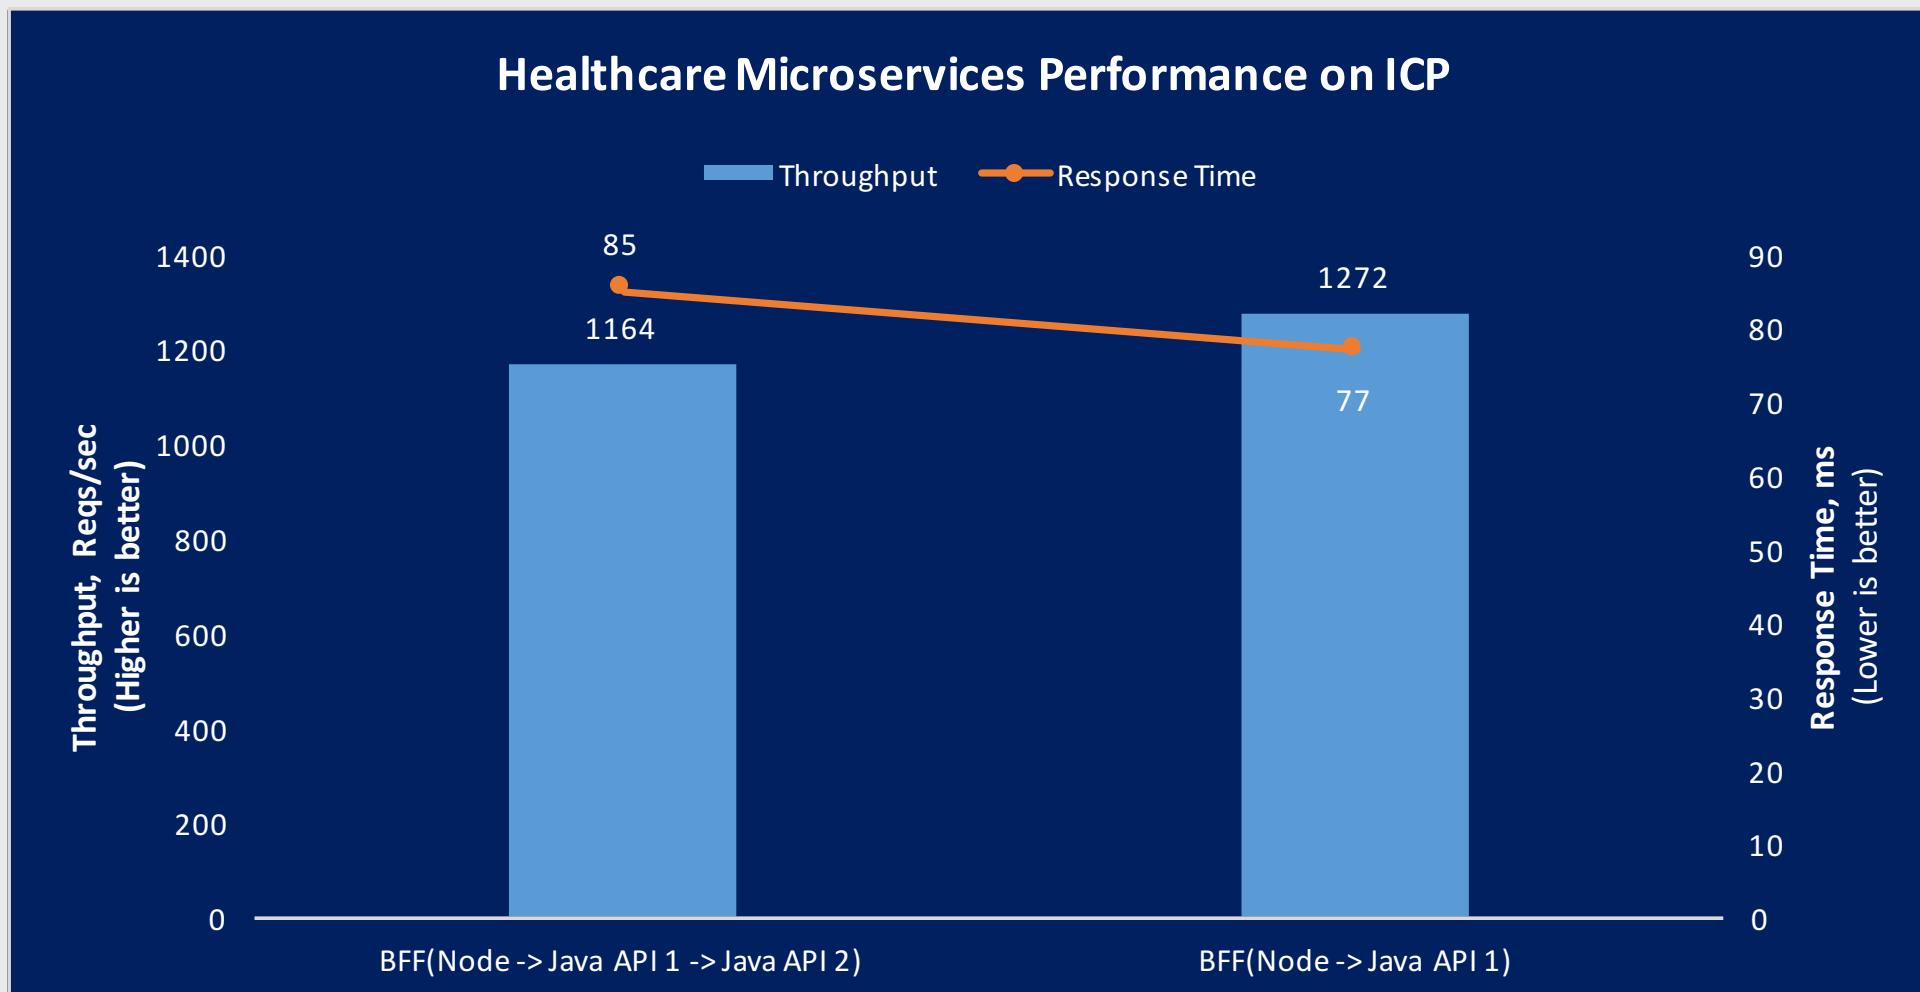
# Healthcare Application Performance on ICP

Comparison of 2 and 3 level BFF service calls

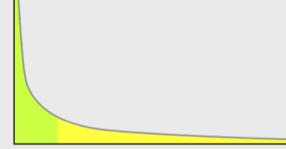
We desire 3 level BFF microservices

Many platforms suffer massive latency when making the jump from 2 level to 3 level

The results of this testing shows only an approximate 10% difference when moving from 2 to 3 layers on ICP



# Healthcare Application Performance on ICP Long-Tail Latency



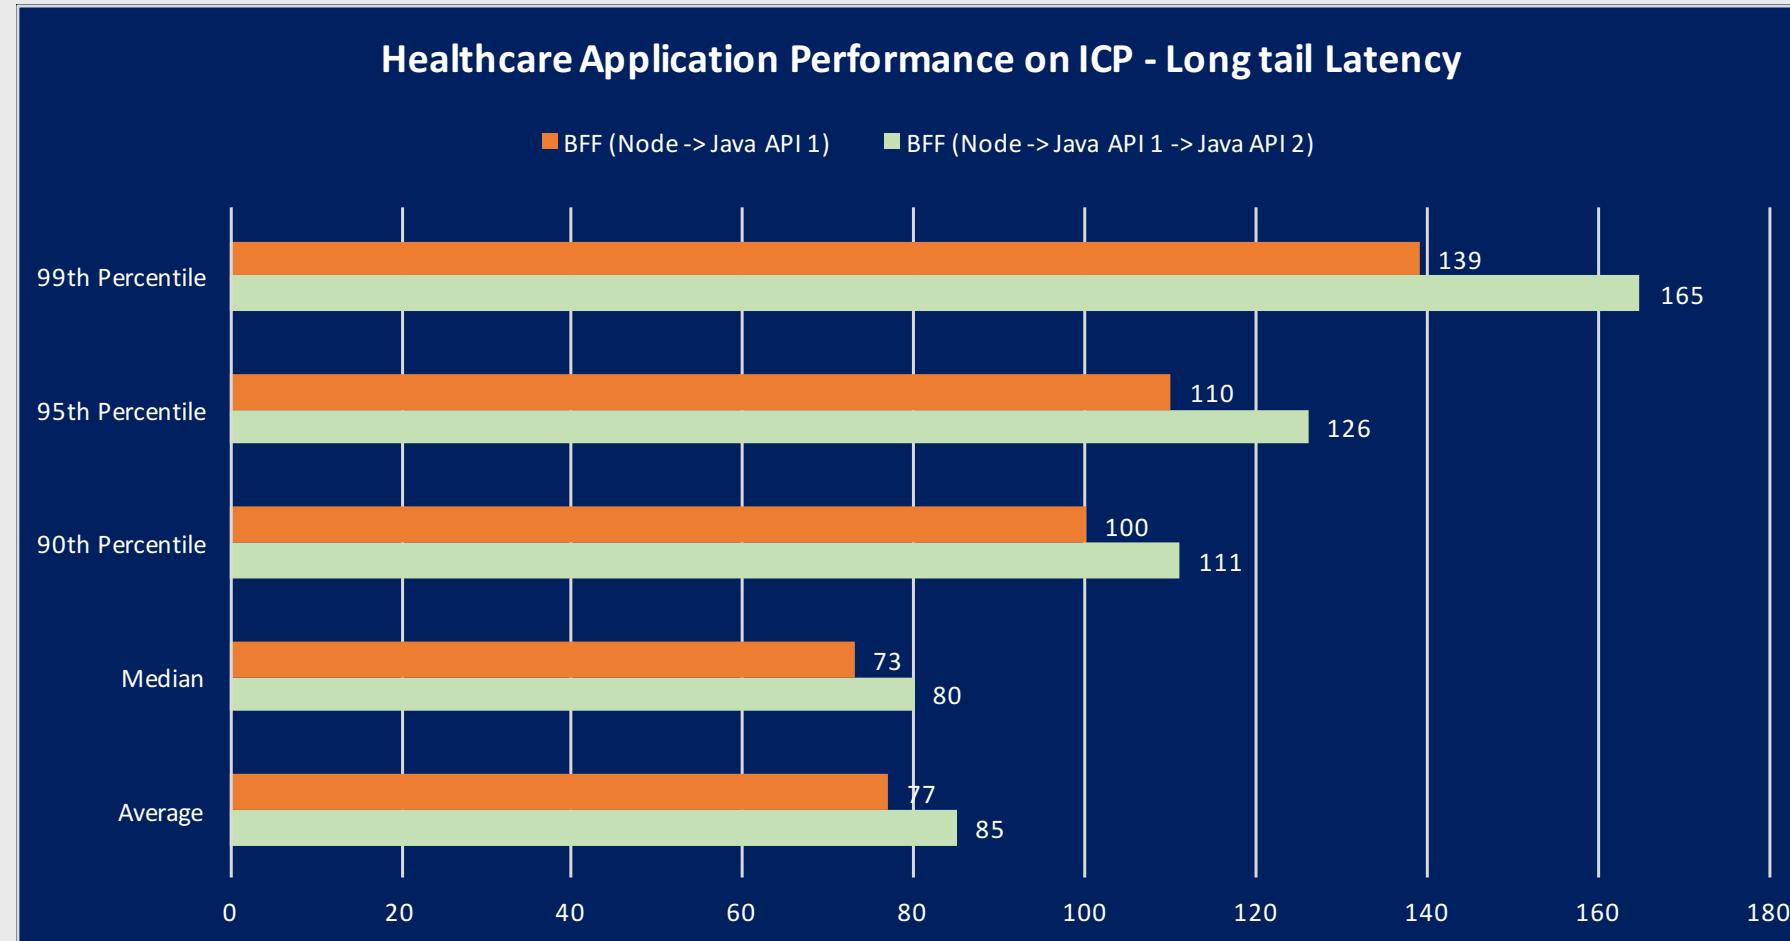
Statistically speaking, the long tail of distributions is the portion of the distribution having a large number of instances or occurrences far from the head (mode or median) of the distribution.

We can consider long-tail latency manifested in the difference between the 90 and 99 percentiles  
Long-Tail latency being high typically results in a negative user experience

It can be a difficult problem to solve

Not uncommon in some instances to see 10X (1000%) difference in this region of a distribution

The test shown here shows 0.45X (45%) for IBM Cloud Private



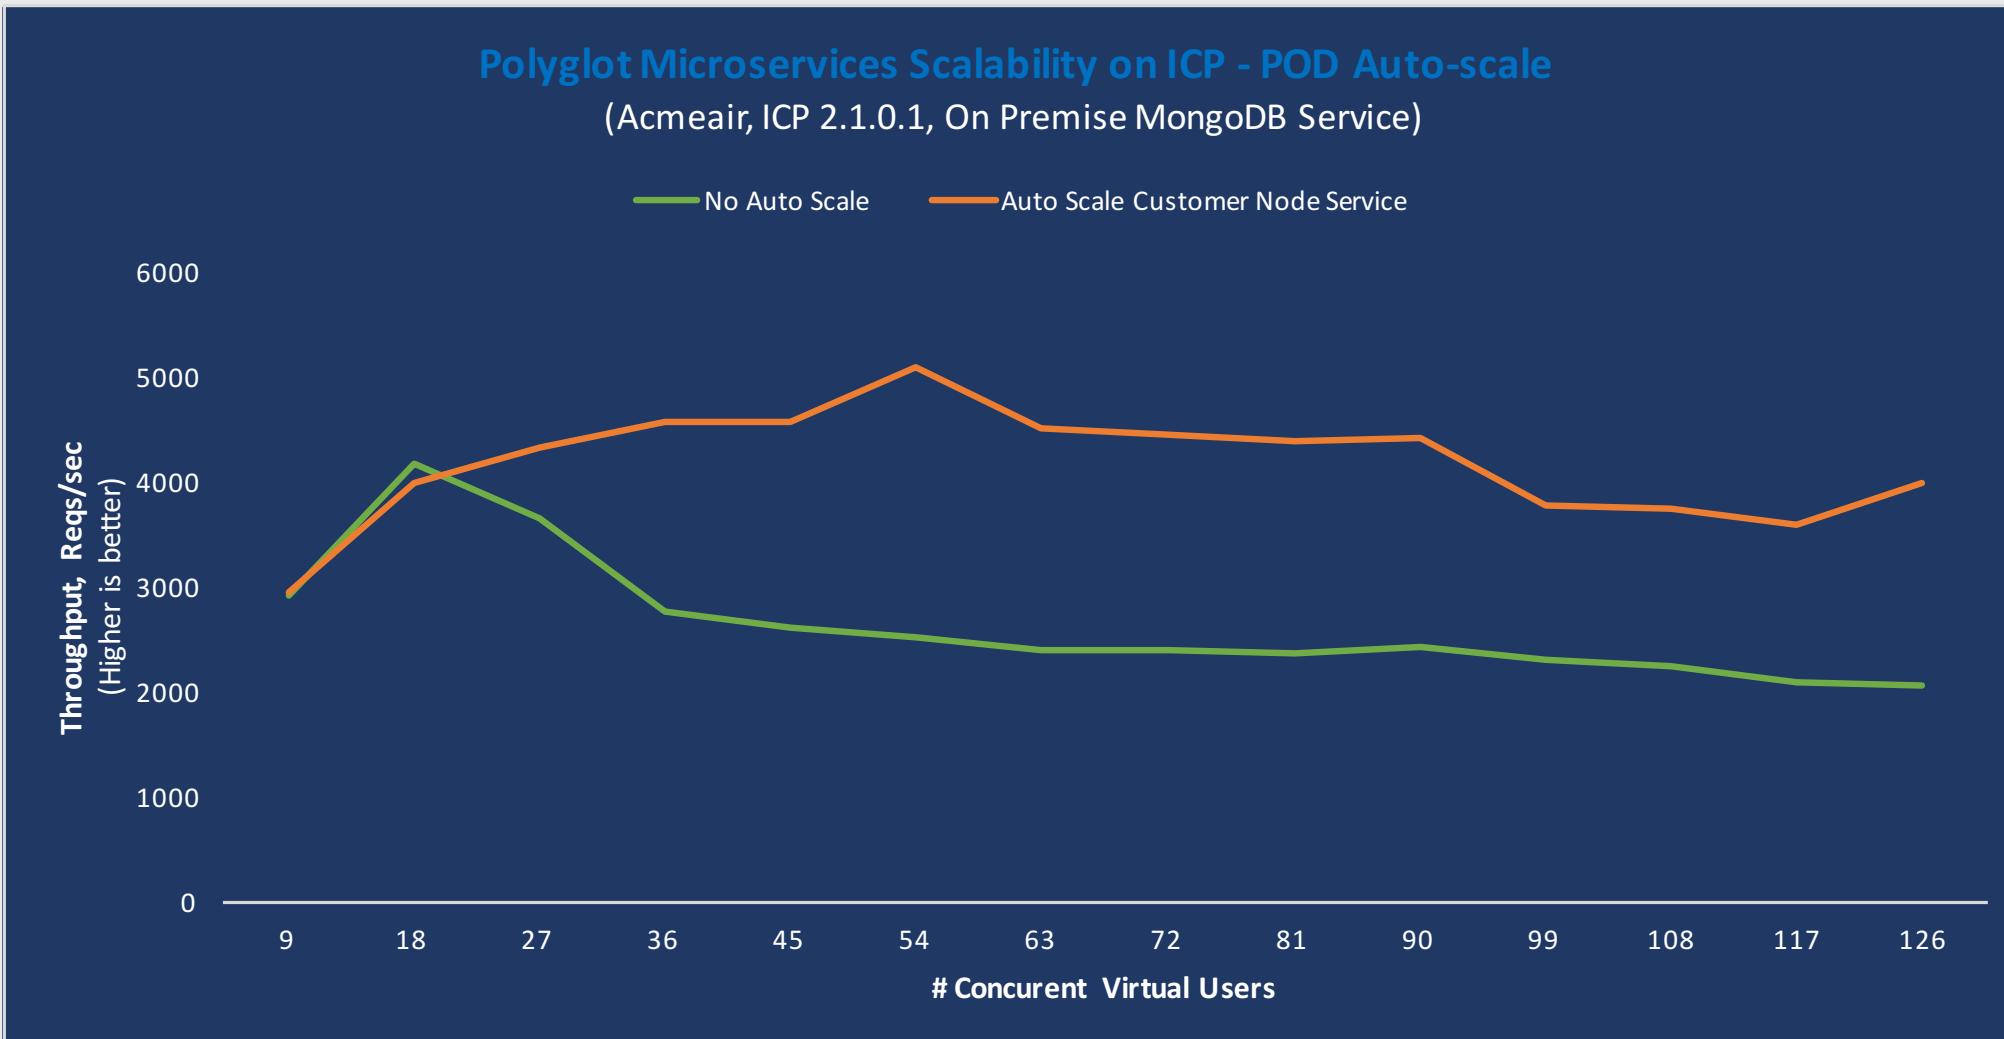
# Polyglot Microservices Performance and Scalability on ICP

The test being run in this example is a node.js application accessing a MongoDB data service  
We are comparing "No" and "Auto" scaling on the node.js service

The services here are deployed within the same cluster

The inability of the application to scale (green line) shows the overall degradation that can occur

These principles still apply as we consider locating the MongoDB externally to the cluster



# Importance of POD Auto-Scaling on Application Performance

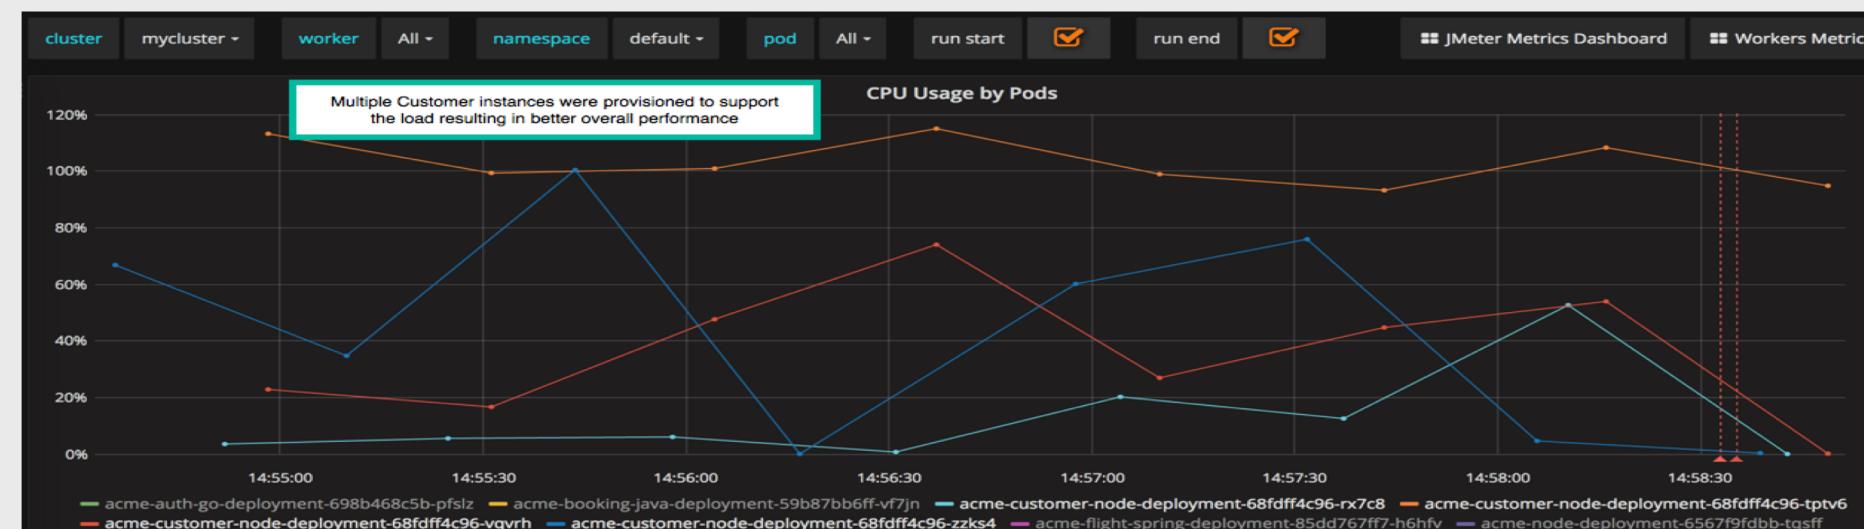
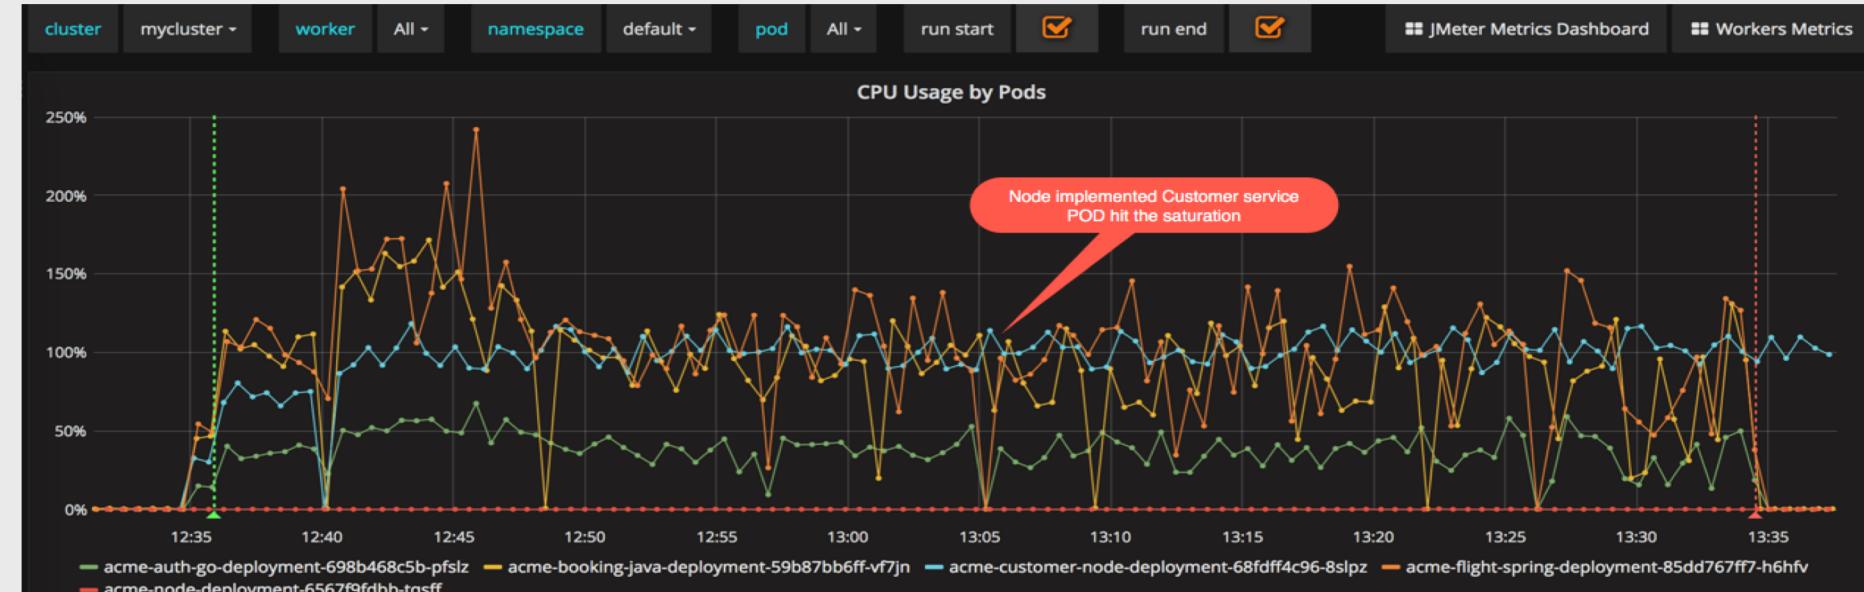
The top chart shows no pod auto-scaling and the bottom chart is with auto-scaling

Top Chart: Blue line shows a pod where the customer service is deployed

Top Chart: Performance is for the most part pegged and demonstrates typical node.js behavior

Top Chart: Unable to push beyond ~100 threads

Bottom Chart: Much better performance as the app scales to match demand

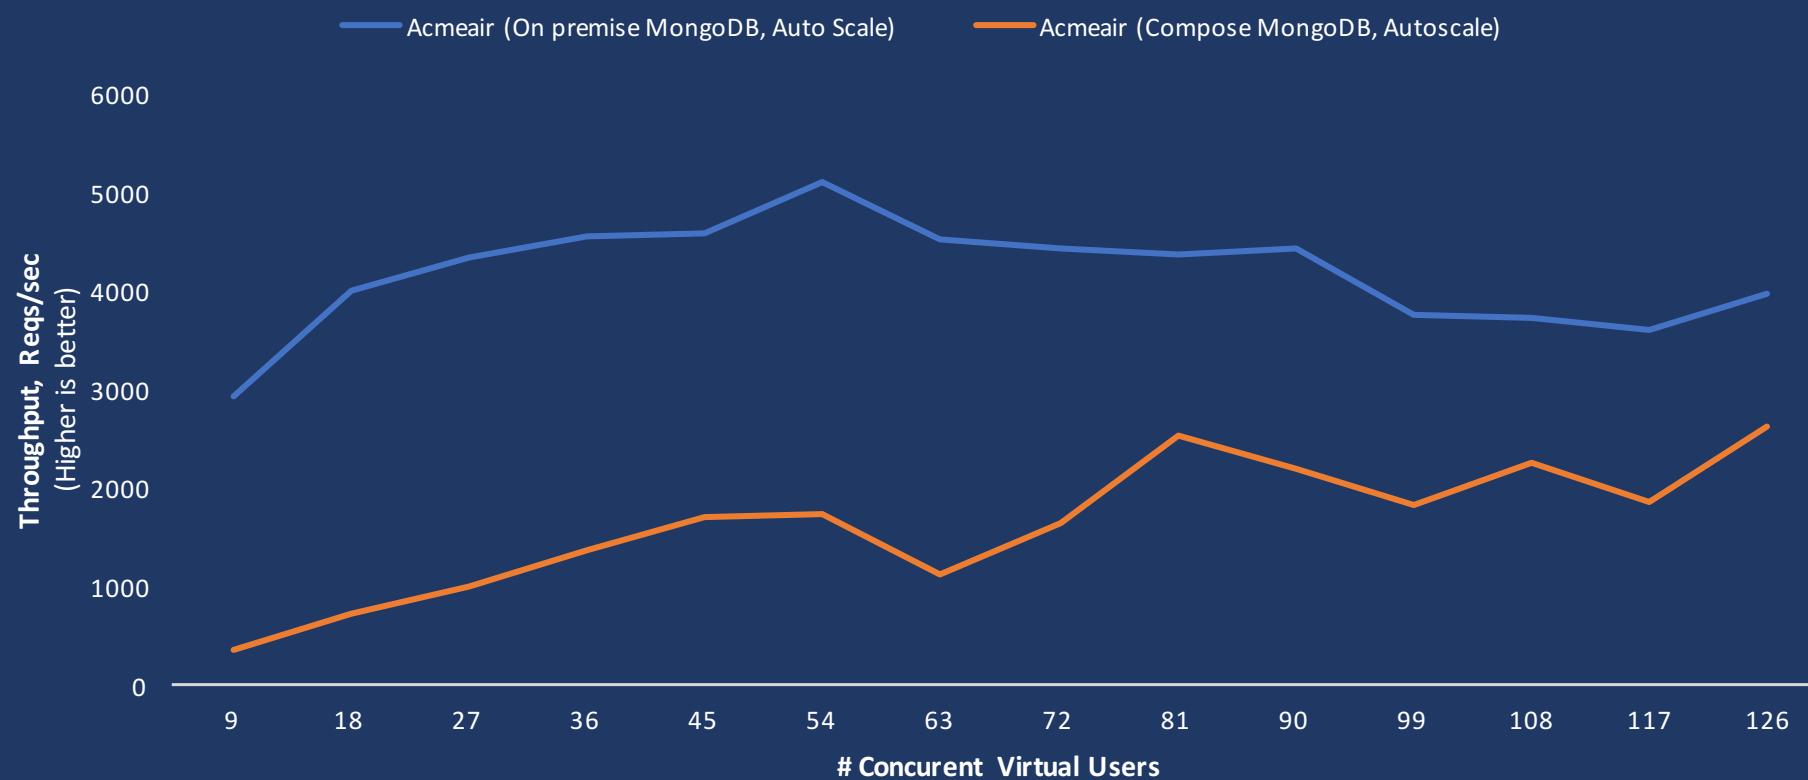


# Compare On Premise v. Public Data Service

Both workloads are shown using auto-scaling

## ICP Performance - Impact of On Premise Data Service vs. Public Data Service

(Acmeair, ICP 2.1.0.1, On Premise MongoDB Service vs. Public MongoDB Service)



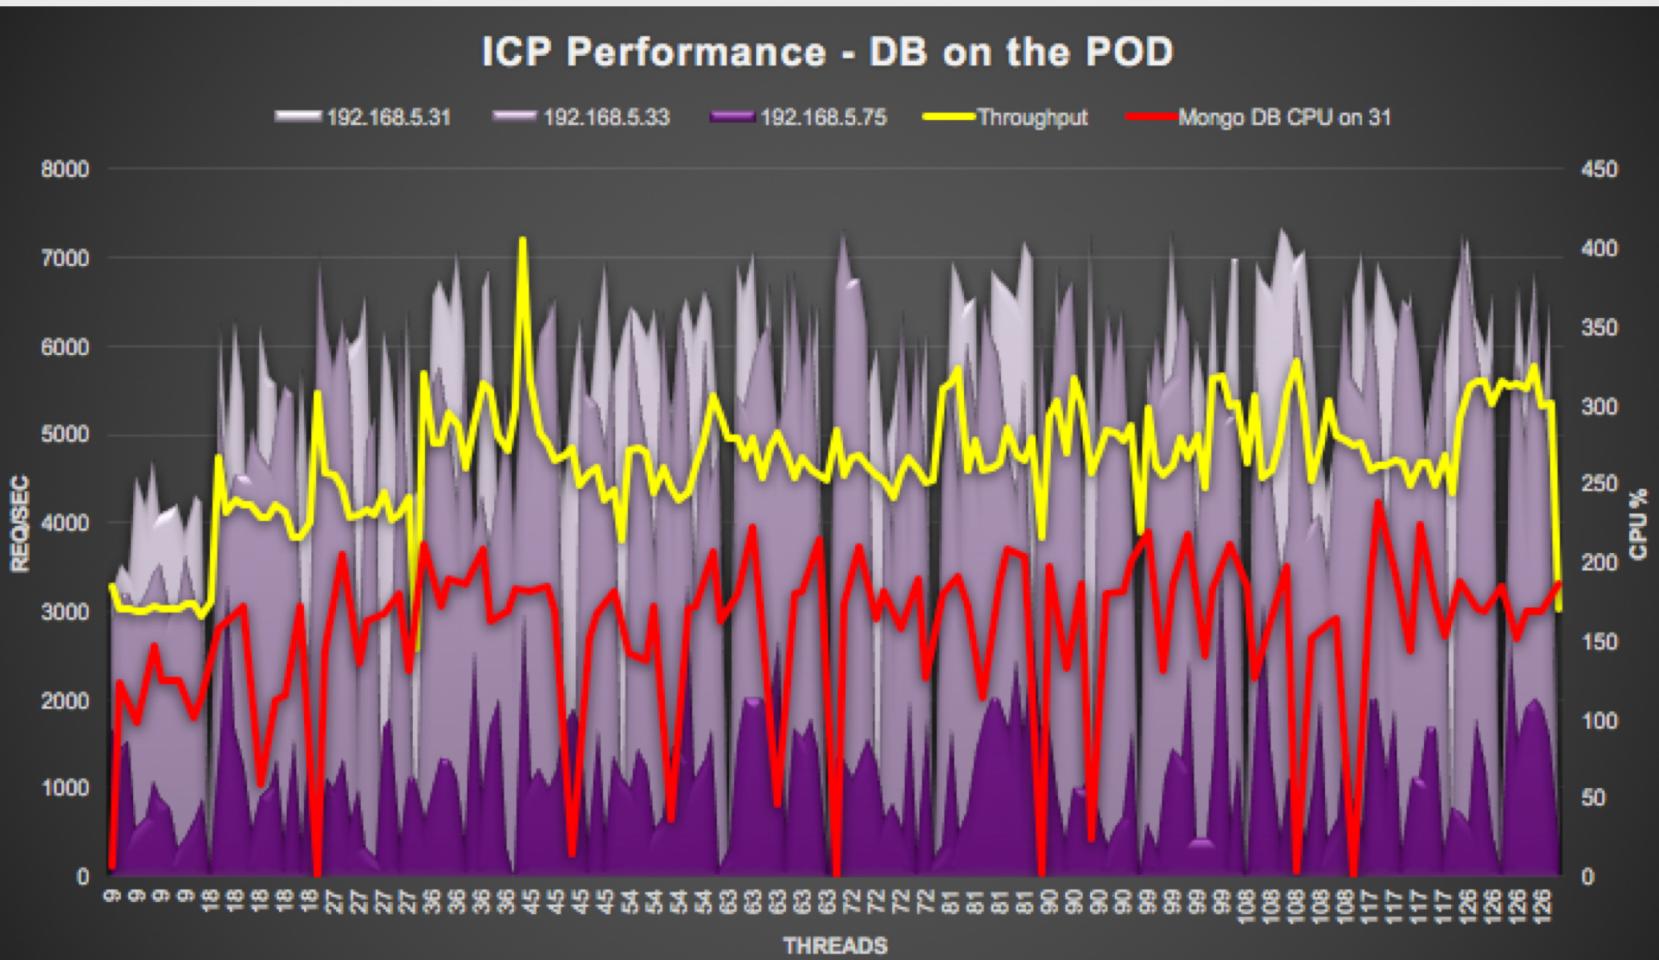
Shows the impact of network latency introduced with the public data service (Compose MongoDB)

We see linear scaling in the public scenario with lower performance throughout as compared to the on premise configuration

# ICP Performance: MongoDB on ICP Cluster

Consider polyglot microservices with MongoDB deployed via Helm locally into the same cluster

Placement of the DB is crucial for scalability of the overall application



The Yellow line shows the throughput of the polyglot application as load is increased

The Red line shows the associated MongoDB pod CPU activity

The remaining (purple) represents the 3 worker nodes in this cluster (saturates at ~400% CPU)

Maximum performance would occur by separating the workload for MongoDB

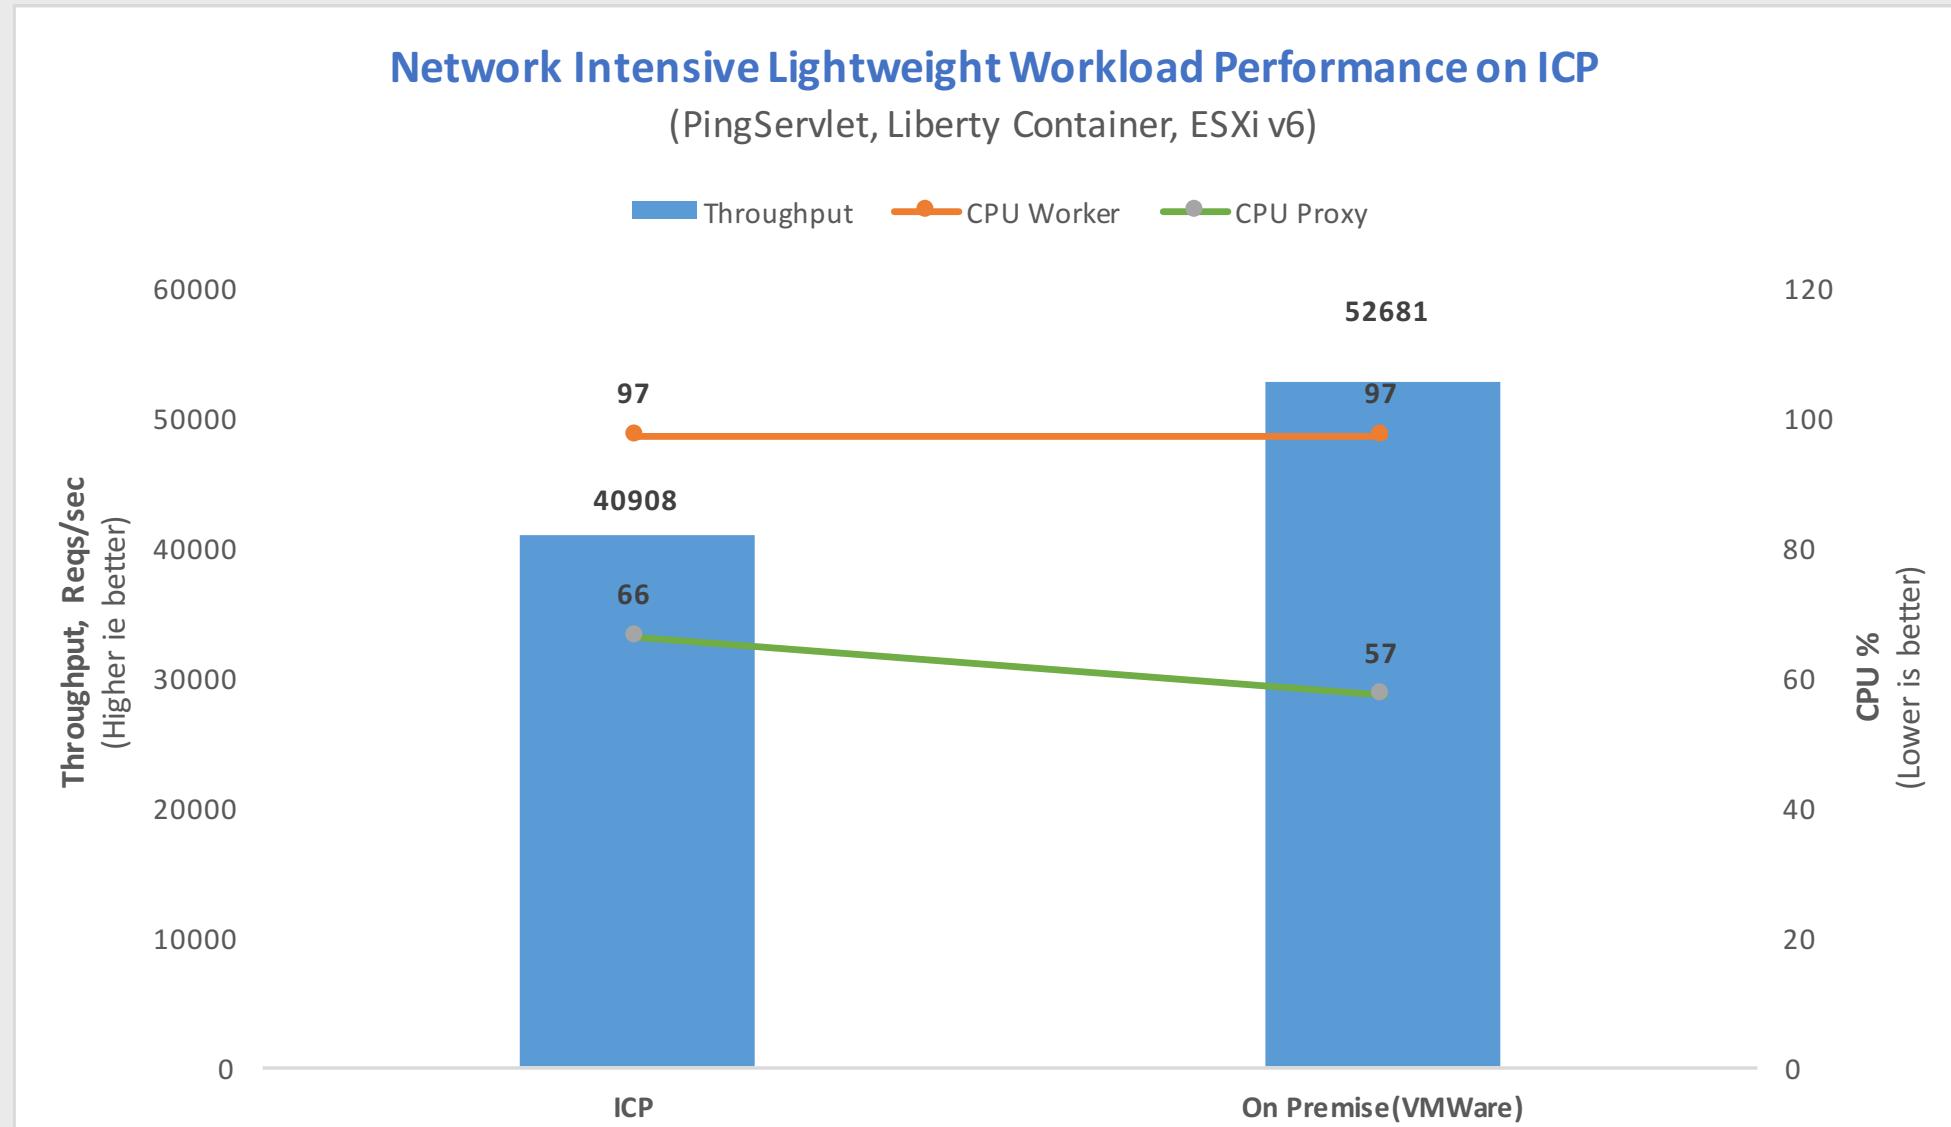
# Network Intensive Application Performance on ICP

Consider the moving of a legacy application that is network intensive to IBM Cloud Private

For network intensive virtualized workloads migrating to ICP we may require 1/3 more compute to support the same transaction load (as network resources transition to CPU)

Docker networking overhead is the source for this load

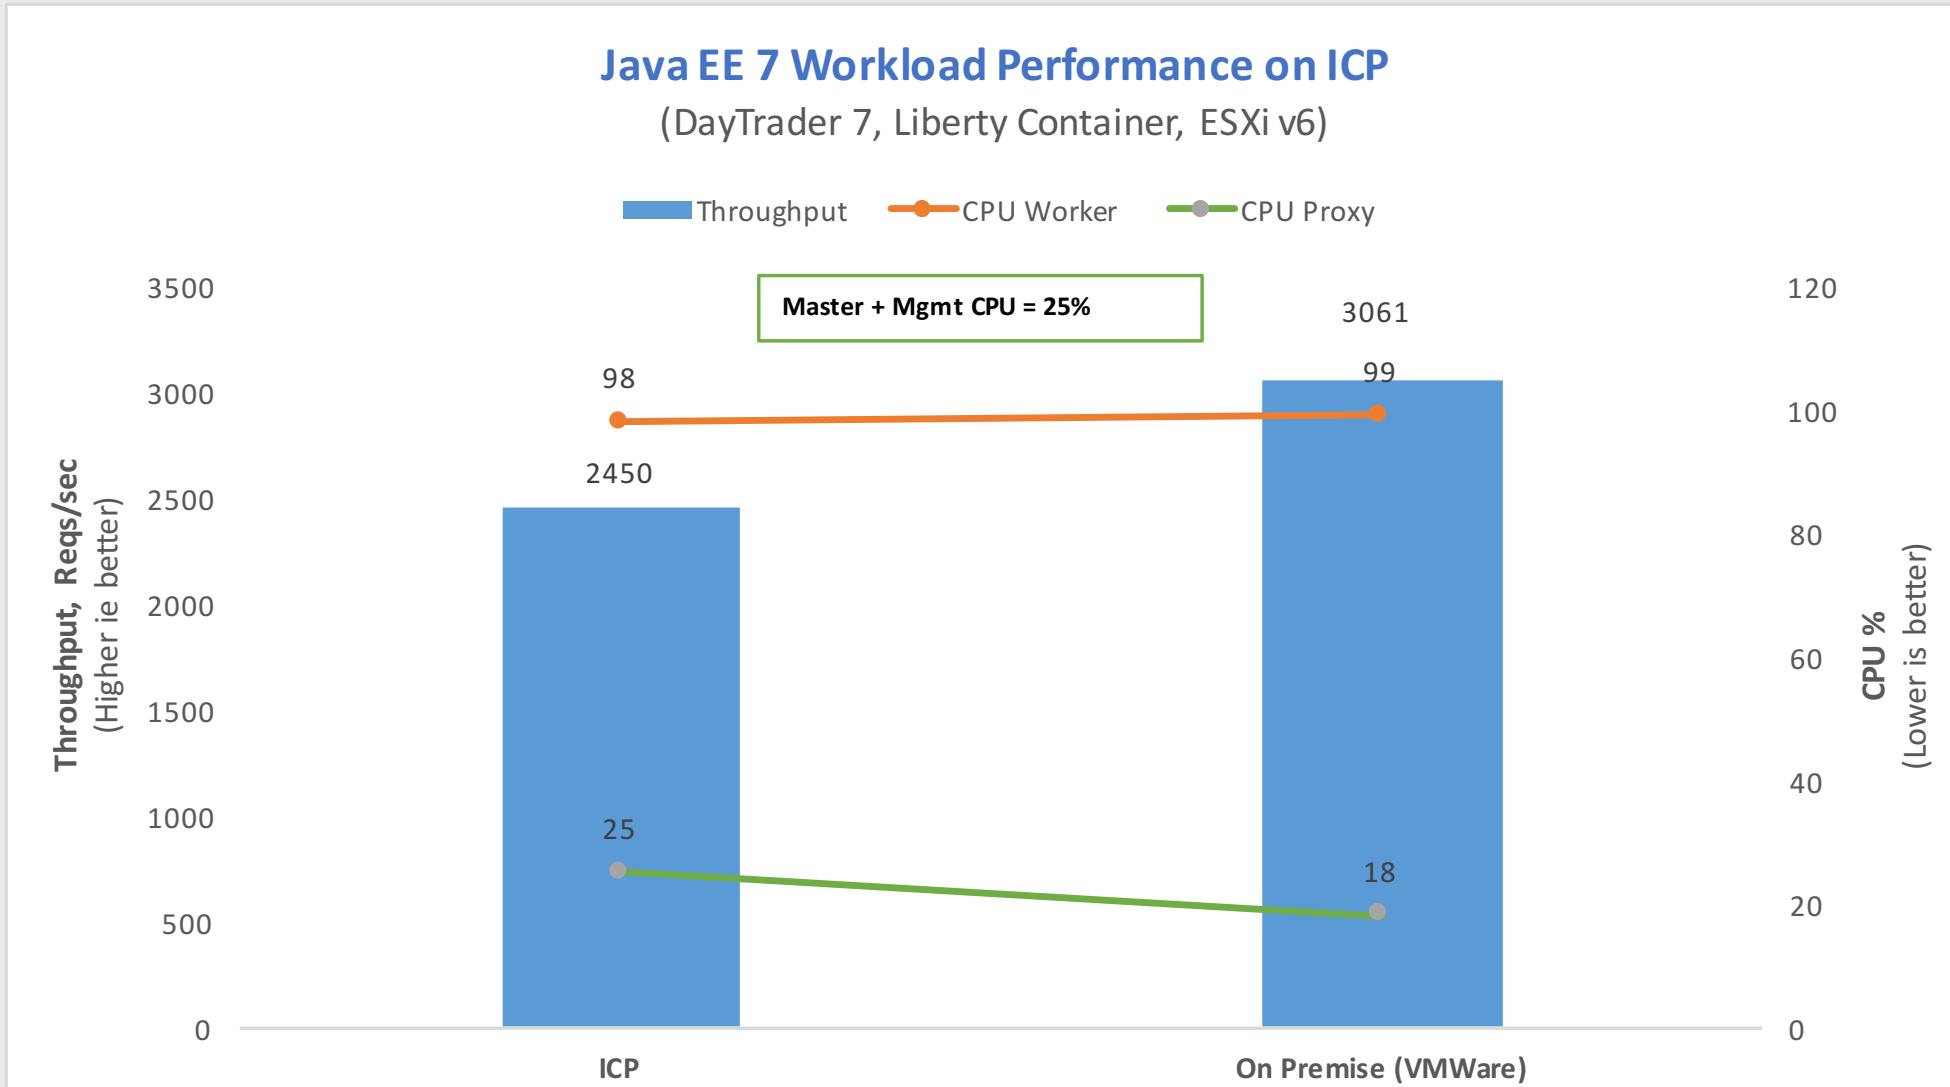
Add an additional 25% capacity for Management / Master nodes and 30% for the running the application



# Java EE Application Performance on ICP

For a Java EE 7 workload migrating from on premise virtualized environment to ICP we require an additional 25% capacity to support similar transaction load (this includes the control plane masters, management, etc.)

This is once again due mostly to Docker networking overhead

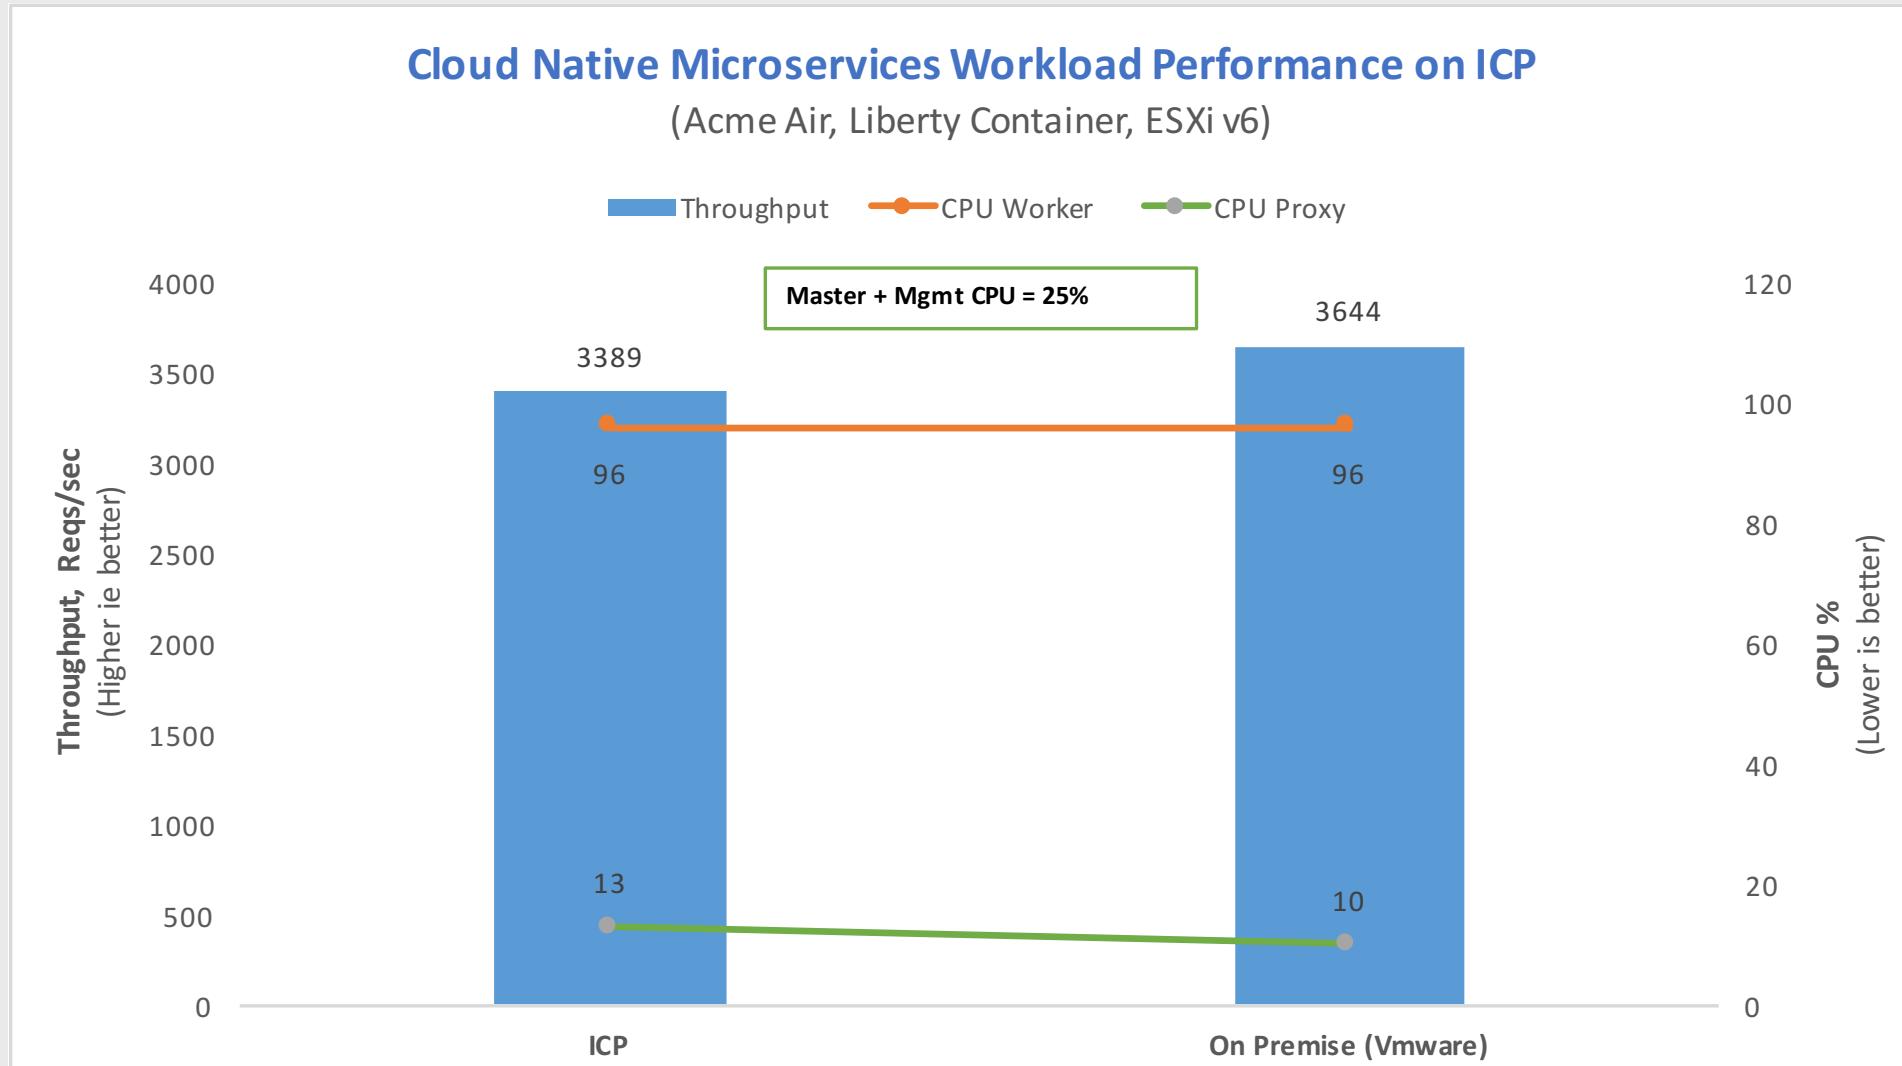


# Cloud Native Application Performance on ICP

For cloud native microservices migrating to ICP from virtualized infrastructure require ~7% additional capacity

Does not include the cluster overhead of ~25% infrastructure increase for the control plane (Master / Management)

Due mostly to Docker networking overhead



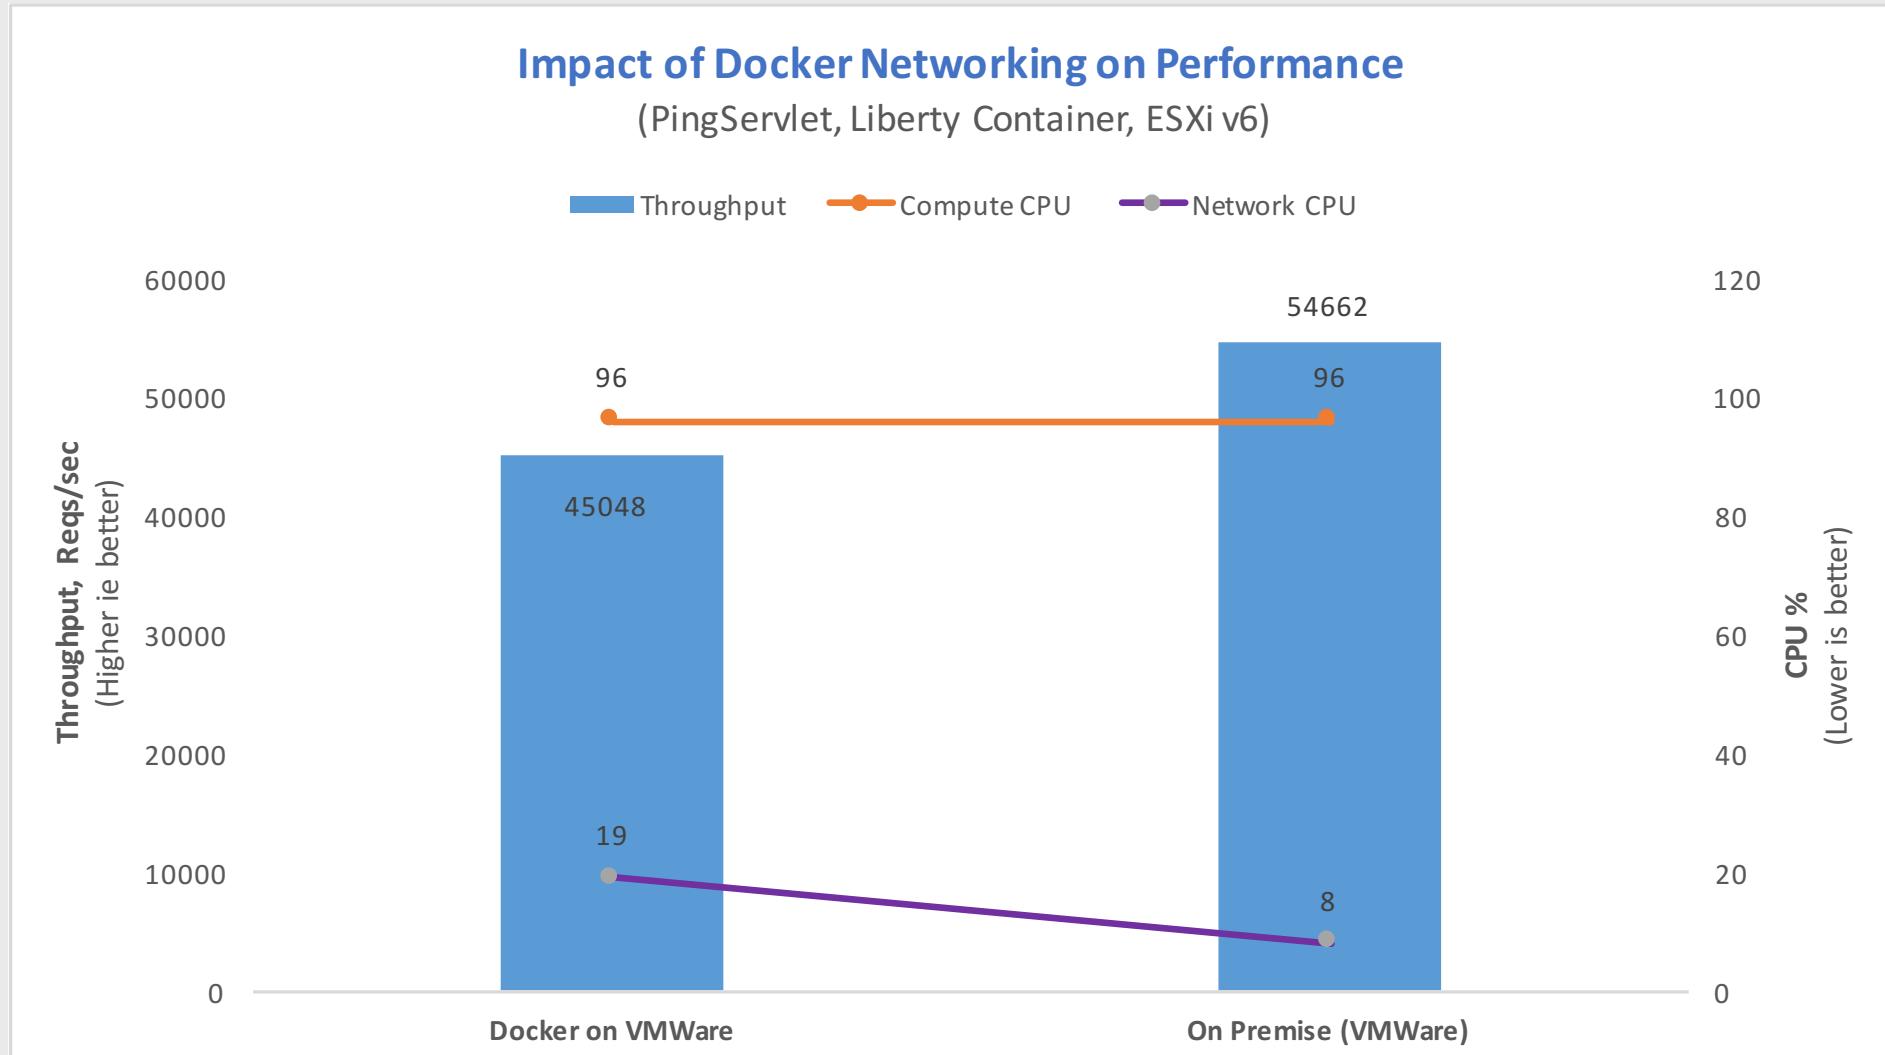
# Reason for Additional Overhead on ICP

Compare Docker on virtualized infrastructure with virtualized alone

We see overhead from Docker networking

Inefficient handling of software and networking transaction are manifested in higher CPU utilization

This behavior is currently being worked from the platform perspective



# Sizing Your Deployment

Choosing a Topology

# Considering Capacity

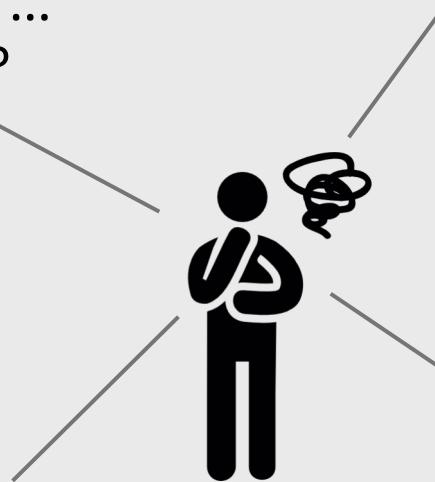
To accurately size your ICP cluster you must consider your workload

What type of applications are you running? Are they microservices ... which frameworks / runtimes ... ?

What are the current resource requirements to run your workloads and how are these requirements expected to change?

Are you running any large containers that require large compute / memory footprints?

Do your applications require persistent or external storage? What type of performance is required for this storage? How many instances of persistent volumes of each type are required?



# The Sandbox

## Assumptions

- Resiliency is a secondary consideration
- Maximum pod size is small
- Small amount of workload
- Vulnerability is not a concern

Role	Count	vCPU / Core	Memory (GB)	Storage (GB)
Master, Boot, Management, etcd	1	4	8	250
Proxy	1	2	4	140
Worker	3	4	8	200
NFS	1	2	4	300

# Production Deployment

## Assumptions

- This is one possible starting point
- Storage must be sized as its own exercise
- Etcd should be broken out when worker nodes surpass 50
- The number and specific size of worker nodes depends heavily on the personality of your workload

Role	Count	vCPU / Core	Memory (GB)	Storage (GB)
Boot	1	4	8	150
Master	3 or 5	8	32	300
Etcd	1 - 3	4	16	150
Management	2 or 3	8	32	300
Proxy	3	4	16	250
Vulnerability Advisor	1 - 3	8	32	500
Worker	5+	8	32	250

Since sizing best practices change regularly with versions and updates, please see this article for the most current shirt size guidelines <https://github.com/ibm-cloud-architecture/refarch-privatecloud/blob/master/Sizing.md>

# Boot Node

Externalizing your boot node alleviates some resource constraints during cluster installation

## Boot Node Considerations

- Externalize for deployment of cluster convenience

# Master Node

Only a single Master Node is active at any time

## Master Node Considerations

- At least 3 nodes ensures a quorum can be reached upon failure
- Additional nodes help ensure the ability to recover from failure
- Consider minimum recommendations to be applicable to only non-production environments
- Can Load Balance masters, but only one is the primary master

# etcd Node

Very large clusters will require externalizing etcd to its own node(s)

## **etcd Node Considerations**

- Shared with Calico
- Scale to meet a larger number of small transactions

# Management Node

Your production clusters  
should use independent  
management nodes

## Management Node Considerations

- Larger clusters with more workload thus require larger management nodes
- Fewer large nodes will have the same impact as many small nodes

# Proxy Node

Understand your workload to tune and size appropriately

## Proxy Node Considerations

- For compute sizing consider total resource sizing versus the # of nodes
- You can tune your ingress controller to your workload
- Your proxy VIP will point only to a single node at a time
- Consider optionally load balancer to spread workload to your proxy nodes

# Worker Node

Consider YOUR workload to size your worker nodes

## Worker Node Considerations

- If your cluster has a small amount of workload consider increasing the number of worker nodes while decreasing the size of the nodes (for adequate headspace, efficiency, mobility, resiliency)
- If you have large pods your worker nodes should be larger to accommodate workload mobility
- Java workloads typically use 4 x CPU for memory
- Other application frameworks may be closer to 2 x CPU = Memory

# Tuning IBM Cloud Private

Best-Practices for Tuning the Cluster

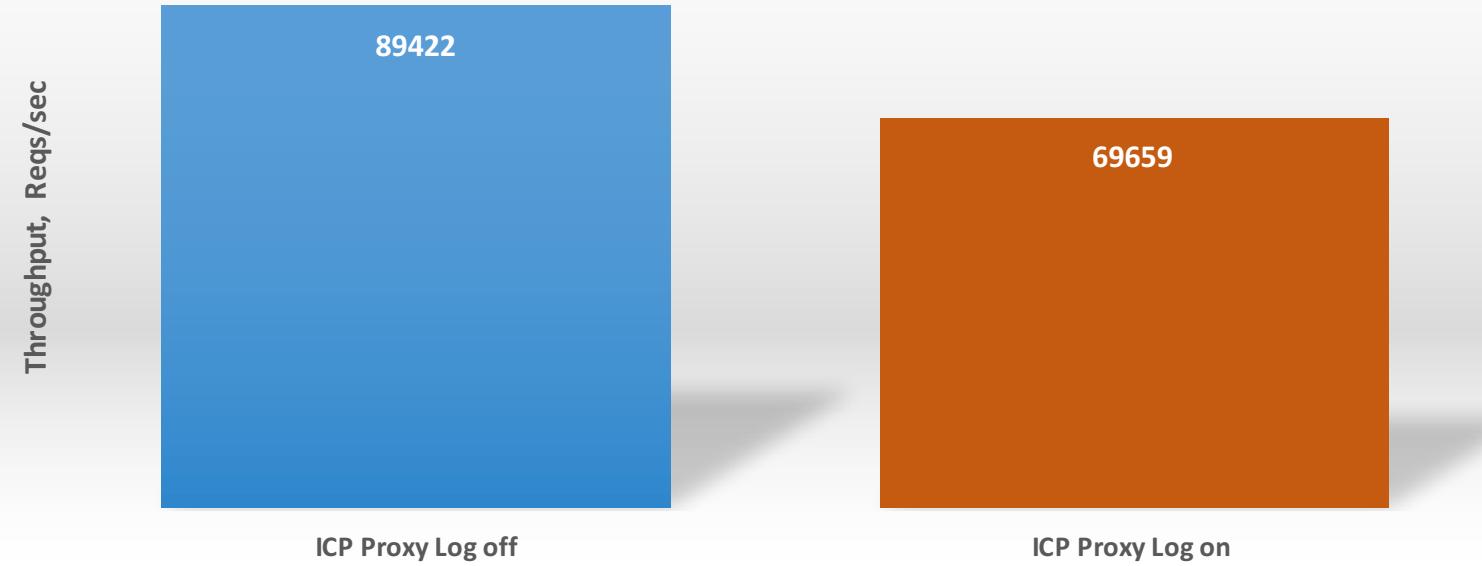
# ICP Tuning: The Impact of ICP Access Logging

Performance is impacted by  
Proxy logging (comparing blue  
and orange bars in both graphs)

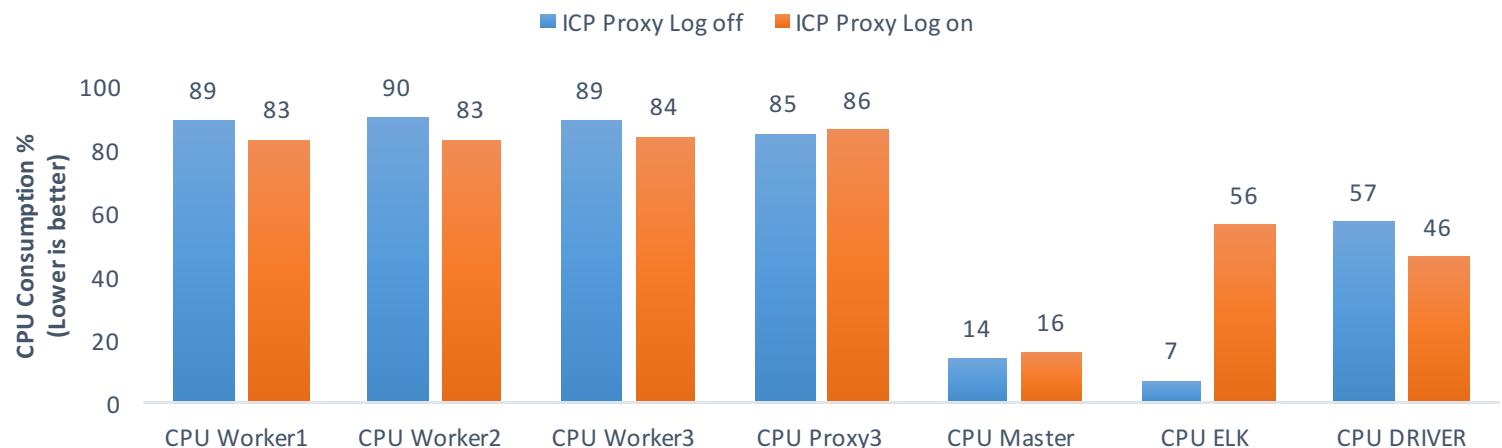
Network intensive loads show  
this can accumulate up to 30%  
additional overhead

The overhead is mostly  
attributable to ELK stack activity

**Impact of Proxy Logging on ICP Performance**  
(Keep Alive 256, 8 vCPU Proxy, PingServlet, Liberty 17.0.0.4, ICP 2.1.0.1)



**ICP Logging Performance - ELK Stack Overhead**  
(PingServlet, Liberty 17.0.0.4, Keep Alive 256, 3 Node ICP 2.1.0.1 Cluster)



# ICP Application Logging

Configure log rotate for Docker  
(max-size, max-file)

Without log rotate disk may fill  
and trigger pod eviction (and  
image garbage collection)

Add install can be configured  
in config.yaml

If Docker is installed  
independently

- Update Docker services  
with “`--log-opt max-size=10m --log-opt max-file=10`”
- Add the config to  
`/etc/docker/daemon.json`

```
## The maximum size of the log before it is rolled
docker_log_max_size: 50m
## The maximum number of log files that can be present
docker_log_max_file: 10
```

## Options

The `json-file` logging driver supports the following logging options:

Option	Description	Example value
<code>max-size</code>	The maximum size of the log before it is rolled. A positive integer plus a modifier representing the unit of measure ( <code>k</code> , <code>m</code> , or <code>g</code> ). Defaults to -1 (unlimited).	<code>--log-opt max-size=10m</code>
<code>max-file</code>	The maximum number of log files that can be present. If rolling the logs creates excess files, the oldest file is removed. Only effective when <code>max-size</code> is also set. A positive integer. Defaults to 1.	<code>--log-opt max-file=3</code>

```
{
  "log-driver": "json-file",
  "log-options": {
    "max-size": "10m"
    "max-file": "10"
  }
}
```

# Impact of Proxy Size and Upstream Keep Alive

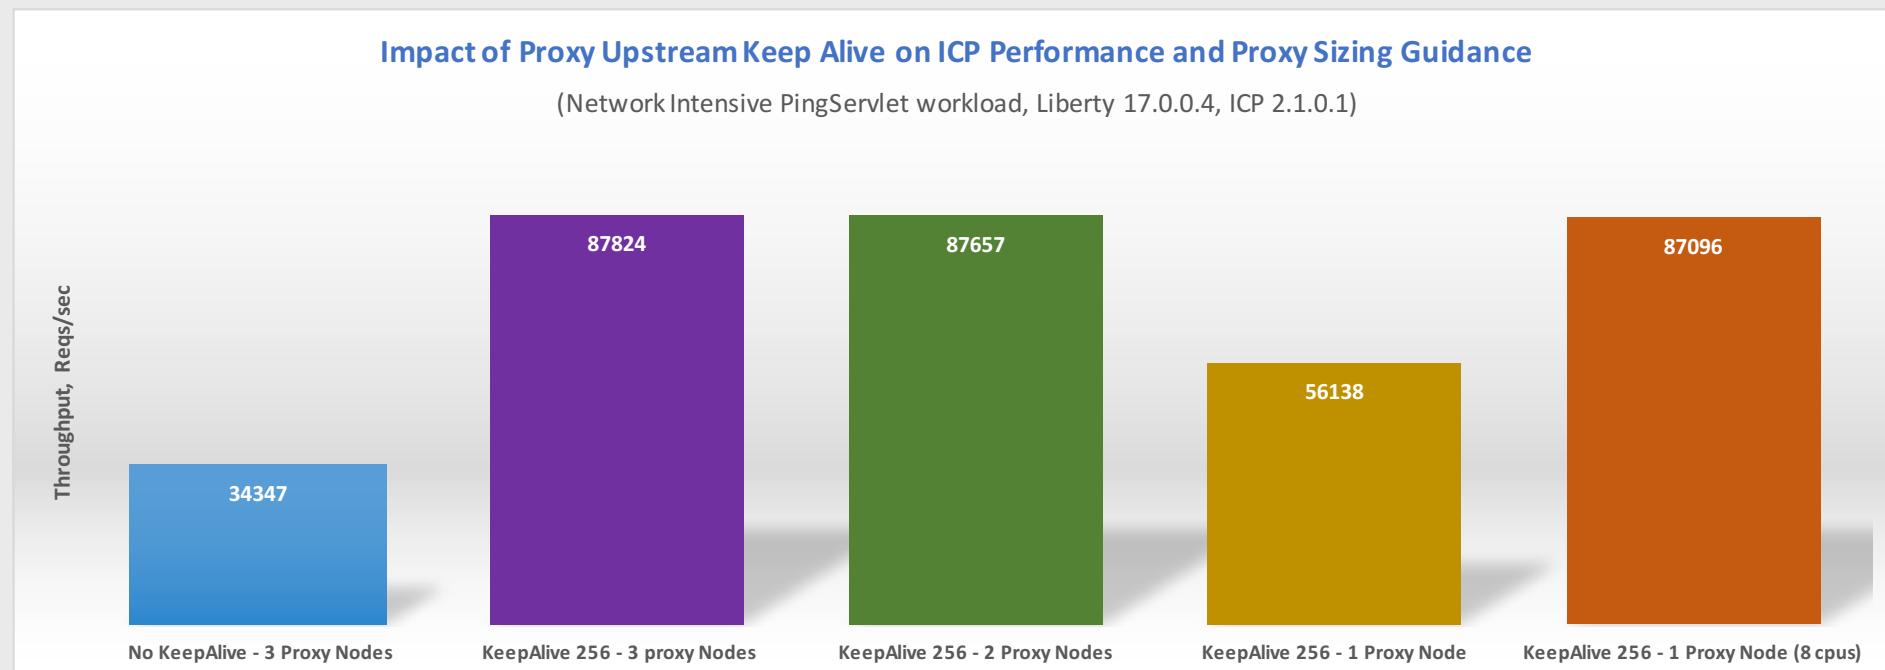
Shows the importance of upstream keep-alive

Proxy node capacity is vital, ensure yours is sufficient

For network intensive workloads, bigger proxy node can be more efficient than multiple smaller proxy nodes

Default keep-alive is 32 – consider tuning this based upon your workload

Note: Keep-alive has been enabled by default since ICP 2.1.0.2



# Ingress Parameters

- Ingress parameters matter →
- Check the configmap
- Update this configmap for tuning performance →
- Among the settings to consider for your workload:
  - Max-worker-connections
  - Worker-processes
- Reference the NGINX / K8s documentation for detail description:  
<https://github.com/kubernetes/ingress-nginx>
- Use initContainer to adjust systemctl default values: kubectl patch ds -n kube-system nginx-ingress-lb-amd64 --patch="\$(cat patch.json)" →

```
spec:  
  containers:  
    - args:  
      - /nginx-ingress-controller  
        default backend service ${POD_NAMESPACE}/default http backend  
      - --configmap=${POD_NAMESPACE}/nginx-load-balancer-conf
```

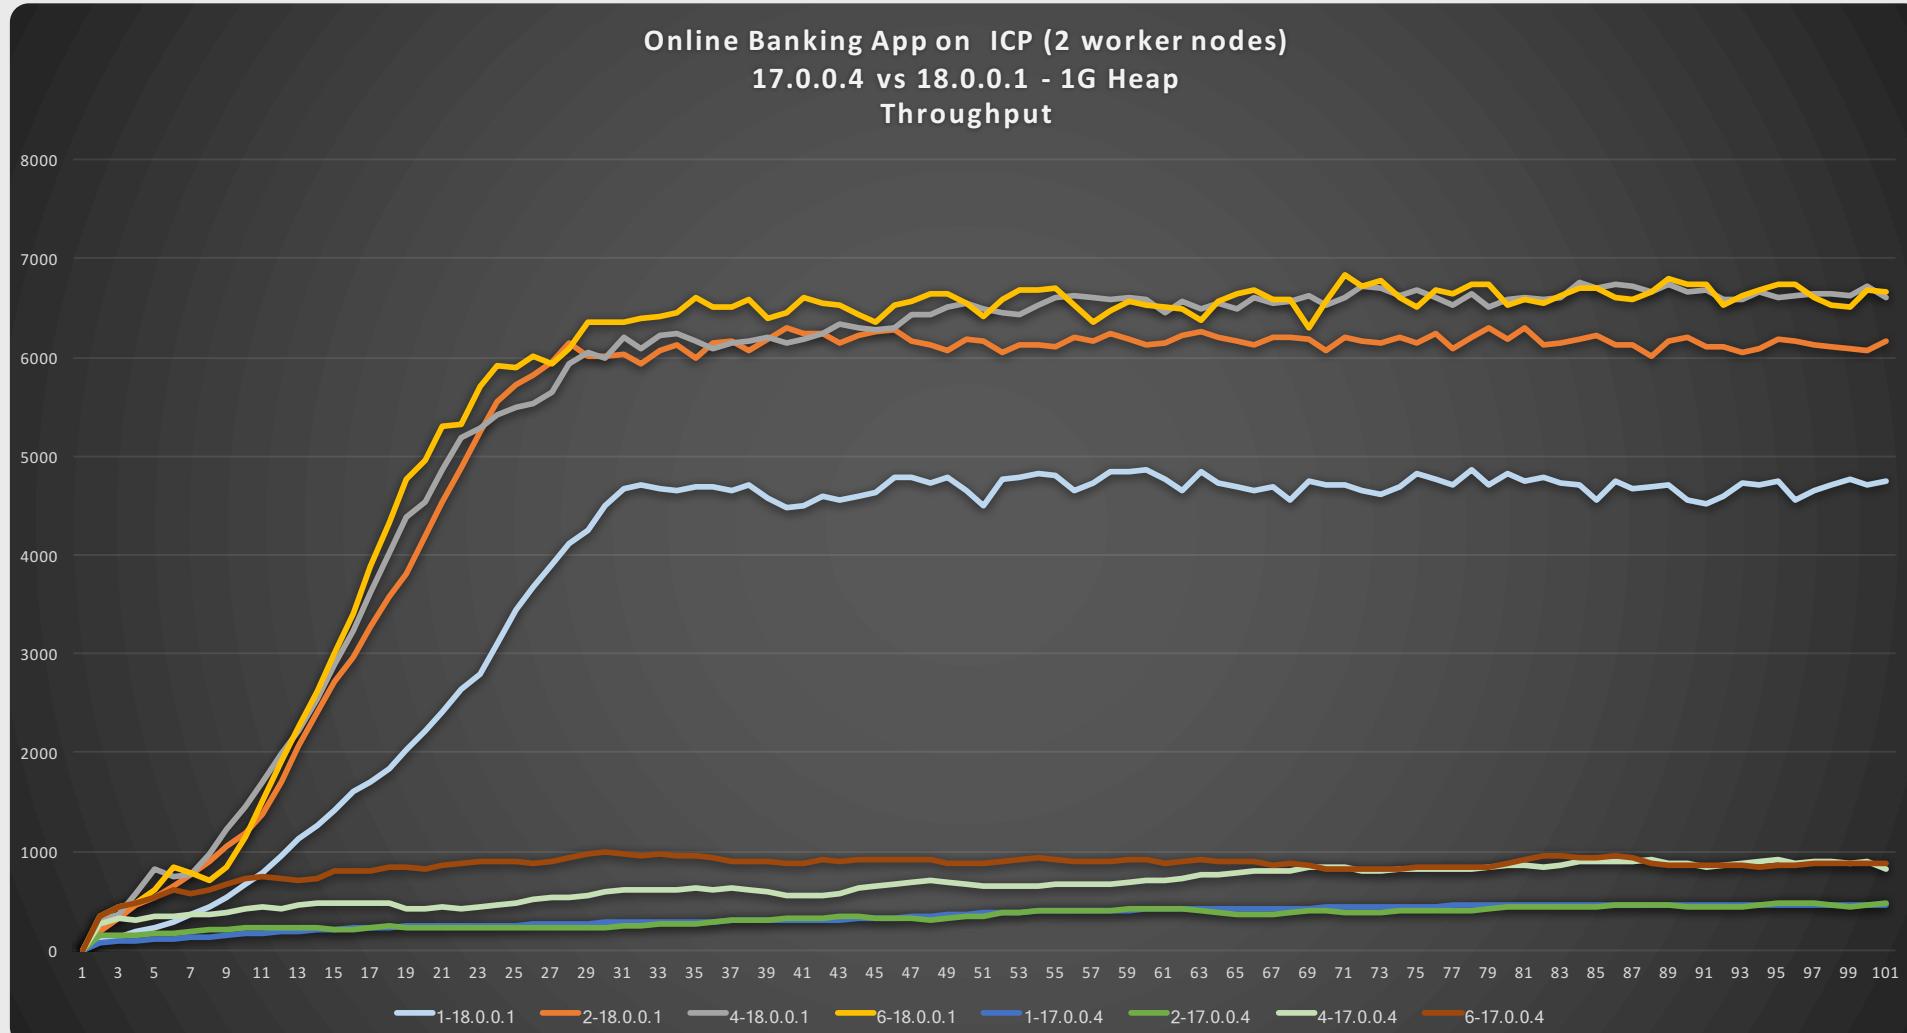
```
root@gyliu-ubuntu-1:/etc/cfc/pods# kubectl get cm -n kube-system nginx-load-balancer-conf -oyaml  
apiVersion: v1  
data:  
  body-size: "0"  
  disable-access-log: "true"  
kind: ConfigMap  
metadata:  
  annotations:  
    kubectl.kubernetes.io/last-applied-configuration: |  
      {"apiVersion": "v1", "data": {"body-size": "0", "disable-access-log": "true"}, "kind": "ConfigMap"}  
  creationTimestamp: 2018-02-23T11:22:31Z  
  name: nginx-load-balancer-conf  
  namespace: kube-system  
  resourceVersion: "1878"  
  selfLink: /api/v1/namespaces/kube-system/configmaps/nginx-load-balancer-conf  
  uid: d6d4c45b-188b-11e8-aae8-fa163e0ee8bb
```

```
{  
  "spec": {  
    "template": {  
      "spec": {  
        "initContainers": [  
          {  
            "name": "sysctl",  
            "image": "alpine:3.6",  
            "securityContext": {  
              "privileged": true  
            },  
            "command": [  
              "sh",  
              "-c",  
              "sysctl -w net.core.somaxconn=32768; sysctl -w net.ipv4.ip_local_port_range=1024 65535"  
            ]  
          }  
        ]  
      }  
    }  
  }  
}
```

# Example: Impact of Liberty Autonomic Threading Algorithm

Java applications running on Liberty with default autonomic threading algorithm may see significant performance degradation

Note: Liberty autonomic algorithm was redesigned in 18.0.0.1 showing a 6X improvement



# **IBM Cloud Private Large Topology**

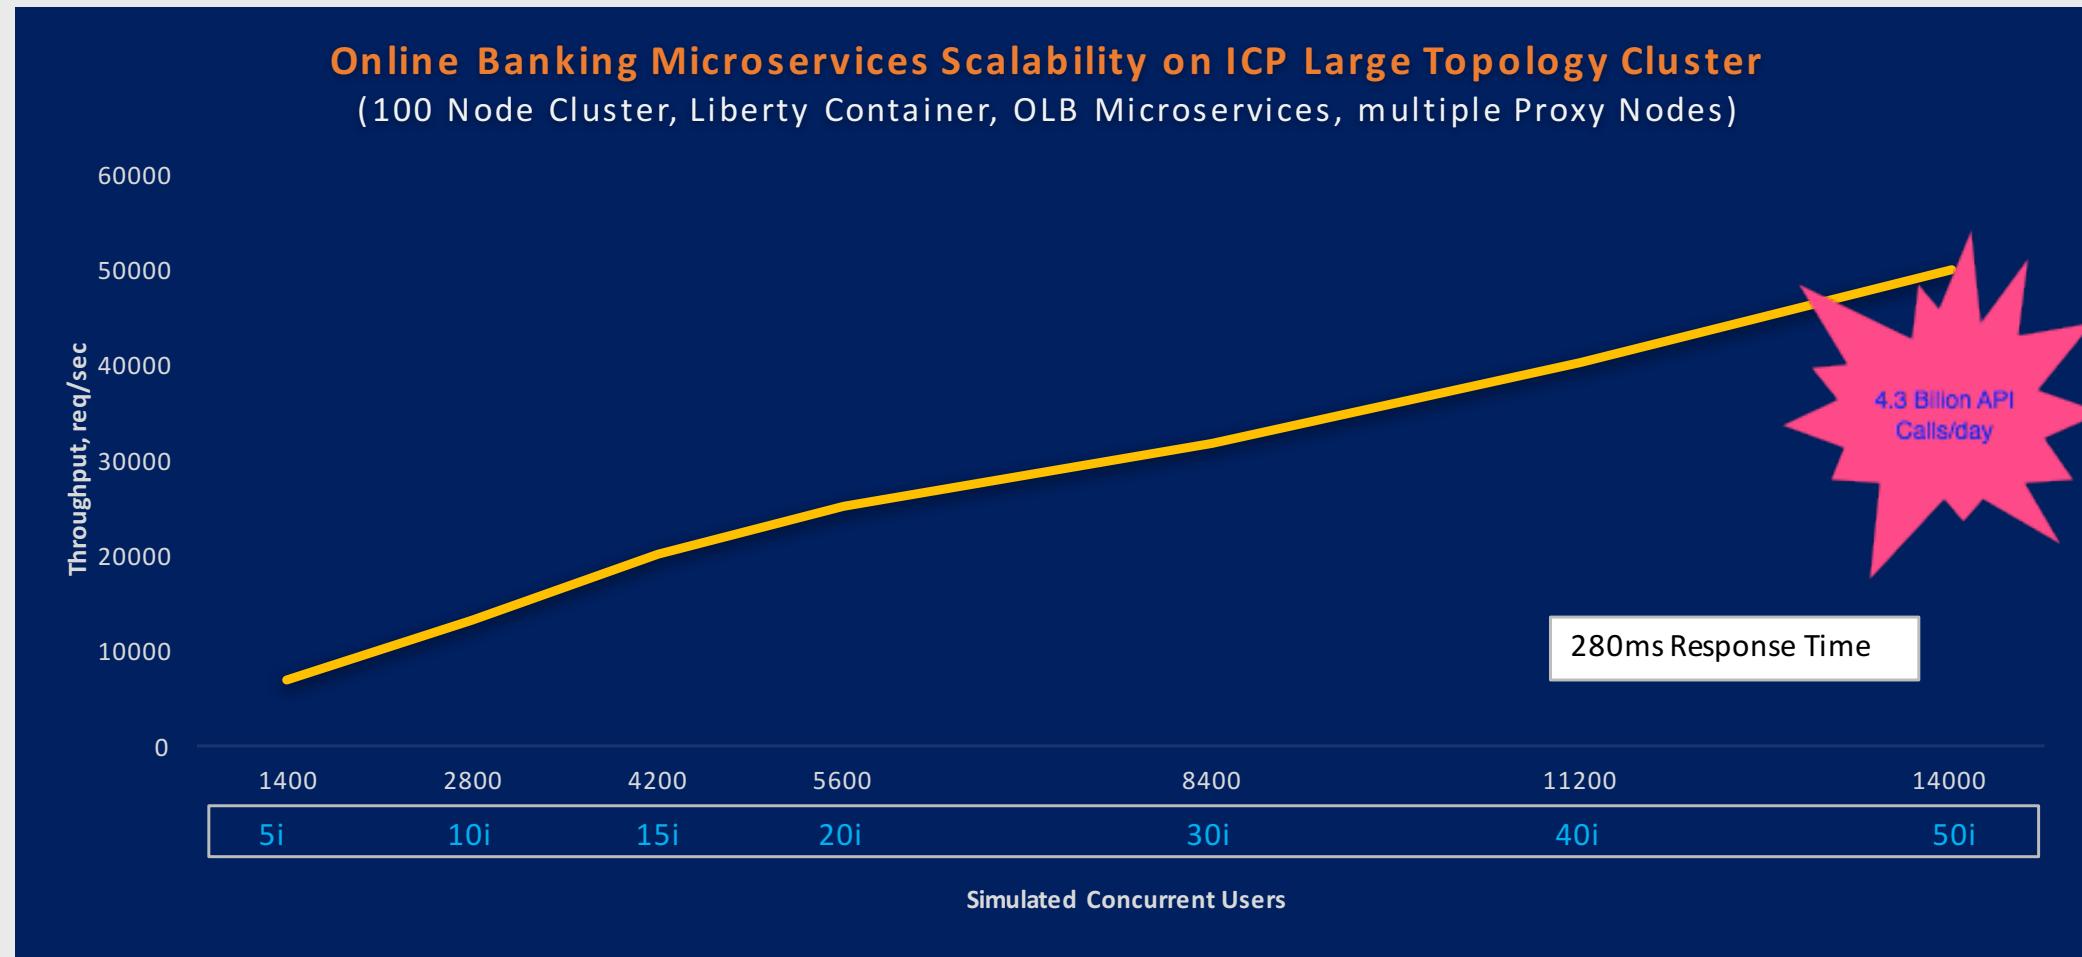
## Performance in Large Topologies

# Large Topology Scalability on ICP Clusters

Online Banking Application

This example is demonstrated on both 100 and 1000 node cluster running 50 instances of the application

The microservices workload scales linearly achieving 4.3 billion API calls per day



# Large Topology Supported Product Limits

ICP Version	Maximum Nodes per Cluster	Total Pods Per Cluster	Pod Spec	Worker Node Spec
2.1.0.1 (Kube 1.8.3)	300	9000	0.1 CPU, 100M Memory	4 Core, 8G Memory
2.1.0.2 (Kube 1.9.1)	500	10000	0.5 CPU, 512M Memory	8 Core, 16G Memory
2.1.0.3	Target 1000	Target 20000	0.5 CPU, 512M Memory	8 Core, 16G Memory

The node sizes are representation not a recommendation with larger clusters typically accommodating larger configurations

The 2.1.0.3 values are unproven or unverified as of the time of creating this presentation

# IBM v. K8s and Supported Product Limits

Kubernetes claims 5000 nodes why can ICP not?

- ICP provides extra enterprise capabilities
- ICP uses Calico as default network overlay
- ICP Calico uses node-to-node mesh to configure peering between all Calico nodes
- Calico load increases etcd load and increases greatly for very large clusters (1000+ nodes)  
  \*\*etcd may be broken out in its own node starting in 2.1.0.3
- Node-to-node mesh stops working if there are more than 700 nodes in a cluster

Note: It is recommended to configure route reflectors for BGP daemons in all cases starting scales approaching 1000 nodes

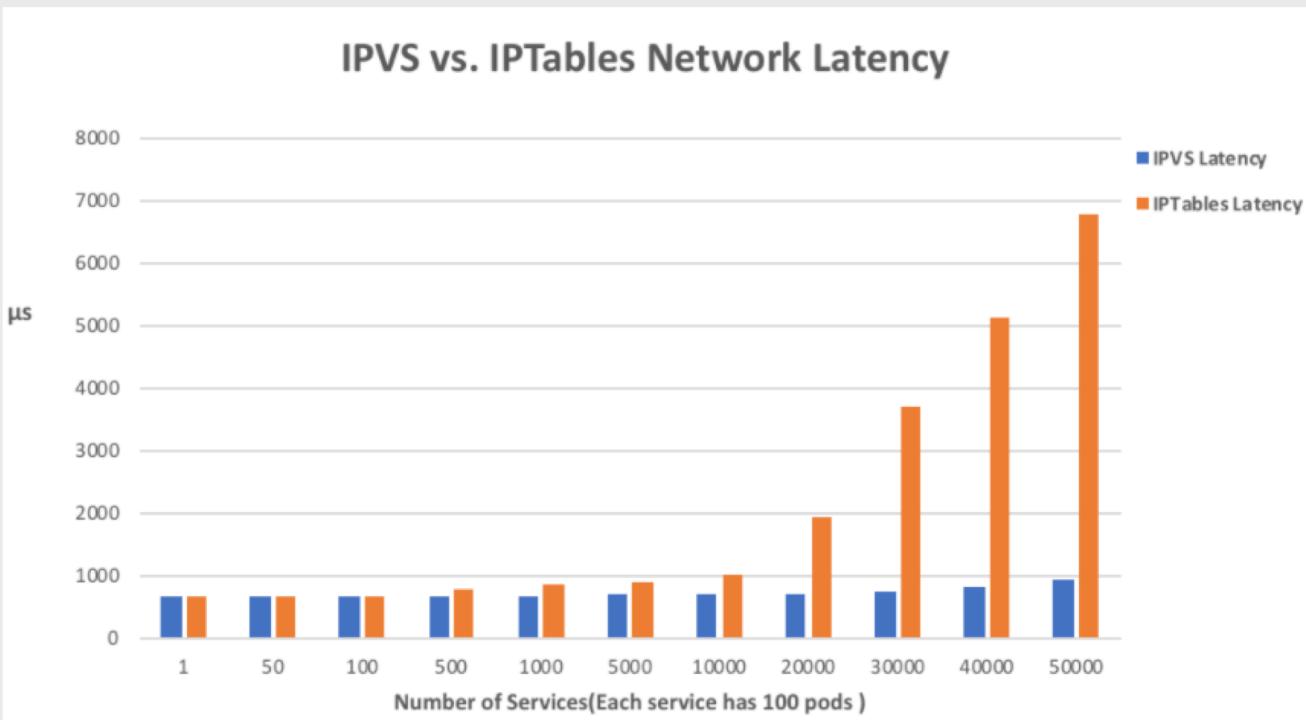
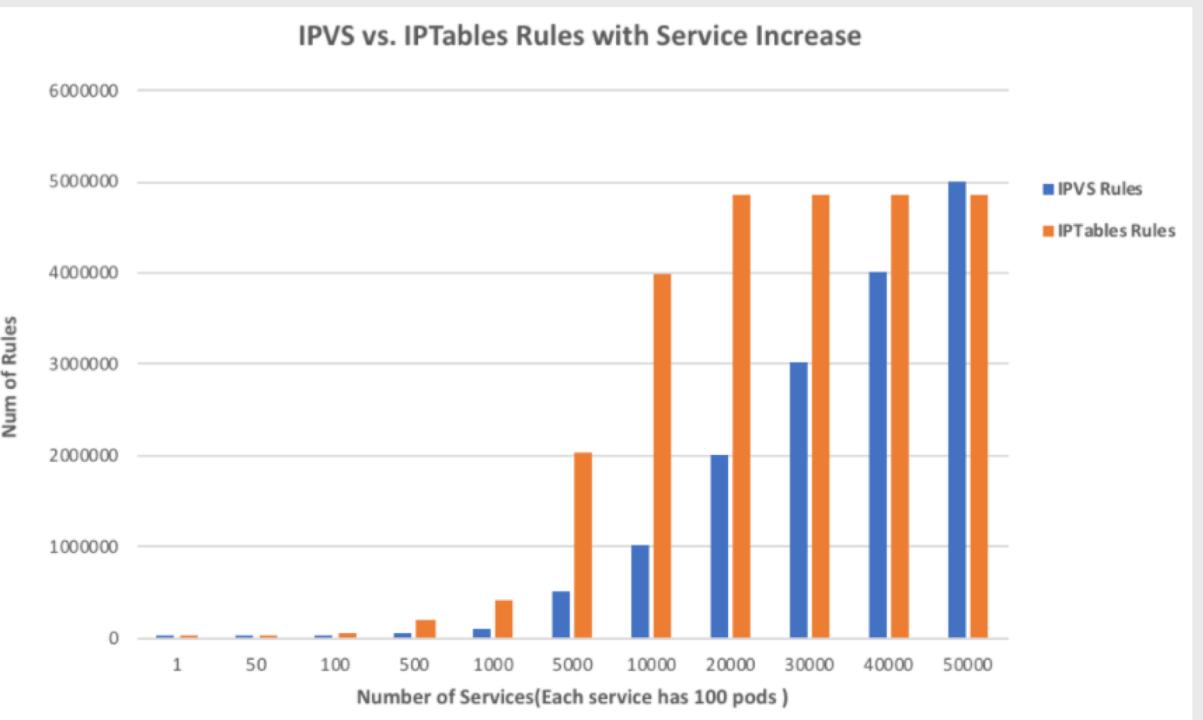
# Kube-proxy: IPVS (Beta K8s 1.10, ICP 3.1.0.3)

We have three choices for kube-proxy mode:  
namespace, iptables (default), ipvs (beta for 2.1.0.3)

Issues with iptables as a load balancers:

- Not incremental: copy all rules, make changes, save back
- Locked during update
- Time to add one rule when there are 5k services (40k rules): 11 minutes increases to 5 hours for 20k services

```
## Kubernetes Settings
# kube_apiserver_extra_args: []
# kube_controller_manager_extra_args: []
kube_proxy_extra_args: ["--feature-gates=SupportIPVSProxyMode=true", "--proxy-mode=ipvs"]
```



# In Summary



## Consider your workload volumes

Do you have a high number of transactions? Is the workload variable?

## Consider your workload personality

Are your services network intensive? What are the programming frameworks you are using?

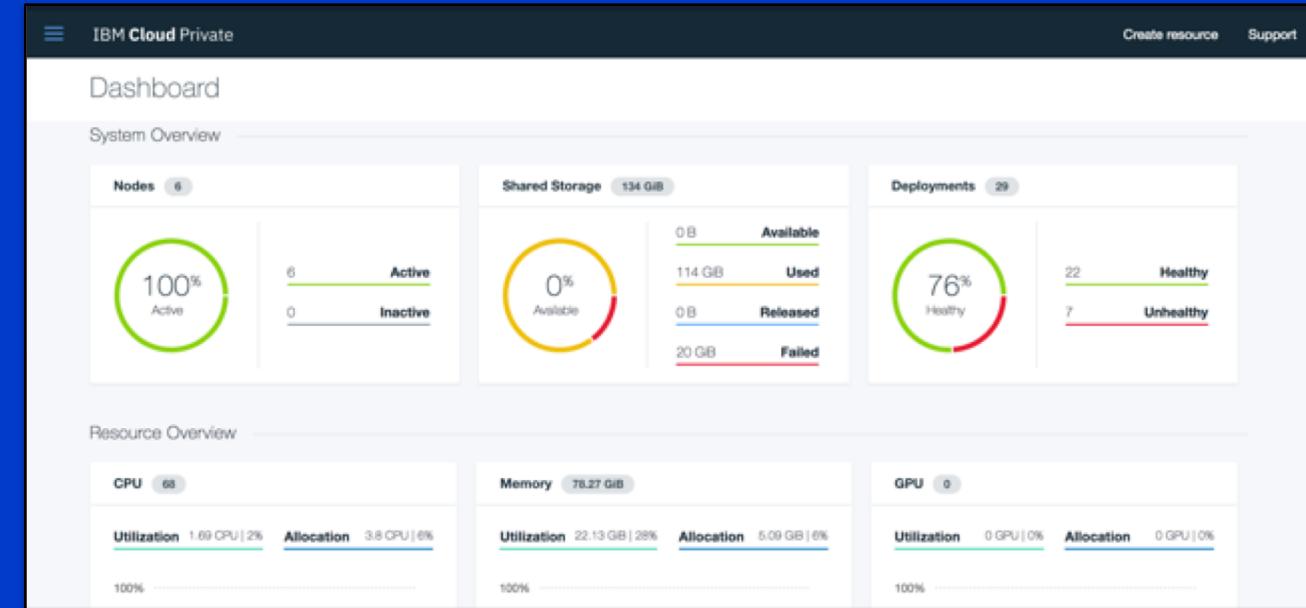
## Consider your performance requirements

How do your applications scale? Horizontal? Vertical? Sharding?

# Try IBM Cloud Private today!

Guided and Proof of  
Technology demos

Free Community Edition!



<http://ibm.biz/ICP-DTE>

